

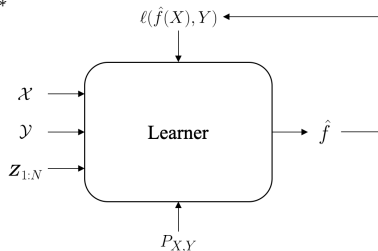
# Topic 1: Learning

# Learning: A Formal Introduction

- ▶ **Domain Set:** An arbitrary set  $\mathcal{X}$  which we may wish to label.
- ▶ **Label Set:** Set of possible labels  $\mathcal{Y} = f^*(\mathcal{X})$ , where  $f^* : \mathcal{X} \rightarrow \mathcal{Y}$  is the true labeling function.
- ▶ **Training Data:** A finite sequence of pairs in  $\mathcal{X} \times \mathcal{Y}$ , denoted as  $\mathbf{z}_{1:N} = (z_1, \dots, z_N) = \left( (x_1, y_1), \dots, (x_N, y_N) \right)^1$
- ▶ **Learner's Output:** A prediction rule  $f : \mathcal{X} \rightarrow \mathcal{Y}$  (also called a predictor, hypothesis, classifier) that predicts the labels of a new domain point from a predictor class  $\mathcal{F}$ . Technically,  $f$  is implemented using an algorithm  $\mathcal{A}(\mathbf{z}_{1:N})$ .
- ▶ **Data Generation Model:** A probability distribution  $P_{X,Y}$  derived from a family of probability distributions  $\mathcal{P}$ , from which the training data  $\mathbf{z}_{1:N}$  is generated.
- ▶ **Measure of Success:** Loss function  $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}_+$  (also called risk, generalized error, true error) that quantifies how bad the prediction rule is, when compared to the true labeling function  $f^*$

**Goal:** Choose the best predictor, i.e.

$$\hat{f} = \underset{f \in \mathcal{F}}{\text{minimize}} \mathbb{E}_{P_{X,Y}} [\ell(f(X), Y)]$$



<sup>1</sup> Sometimes, training data does not come with labels. In such a case, training data is  $\mathbf{x}_{1:N} = (x_1, \dots, x_N)$ .

## Example: Bias Estimation in Coin Tossing

- ▶ Goal: Given a (biased) coin with unknown probability  $\theta$  of turning heads, determine  $\theta$  as accurately as possible.
- ▶ Outcomes of  $N$  coin tosses:  $\mathbf{x}_{1:N} = (x_1, \dots, x_N)$ , where  $x_i = \begin{cases} 1, & \text{if Heads,} \\ 0, & \text{otherwise.} \end{cases}$
- ▶ Given  $\theta$ ,  $N$  and  $\hat{\theta}_N(\cdot)$ , we can partition  $\{0, 1\}^N$ , the set of all  $N$  coin tosses, as

$$\text{Good Data: } G_{N,\epsilon} \triangleq \left\{ \mathbf{x}_{1:N} \in \{0, 1\}^N \mid \left\| \hat{\theta}_N(\mathbf{x}_{1:N}) - \theta \right\| \leq \epsilon \right\}$$

$$\text{Bad Data: } B_{N,\epsilon} \triangleq \left\{ \mathbf{x}_{1:N} \in \{0, 1\}^N \mid \left\| \hat{\theta}_N(\mathbf{x}_{1:N}) - \theta \right\| > \epsilon \right\}$$

### Claim 1

For any true value  $\theta$  of the coin bias, given any  $\epsilon, \delta > 0$ , it suffices to collect  $N \geq \frac{1}{2\epsilon^2} \log \left( \frac{2}{\delta} \right)$  samples to guarantee

$$\mathbb{P}_\theta (G_{N,\epsilon}) = \mathbb{P}_\theta \left( \left\| \hat{\theta}_N(\mathbf{x}_{1:N}) - \theta \right\| \leq \epsilon \right) > 1 - \delta.$$

## Example: Bias Estimation in Coin Tossing (cont...)

Proof of Claim 1:

# Main Essence of Learning

*Our main wish is to learn something about a phenomenon of interest, via observing random samples of a quantity that is relevant to the phenomenon.*

Two basic questions to ask:

- ▶ **Statistical Learning:** How many samples are needed to achieve a given accuracy ( $\epsilon$ ) with a given confidence ( $\delta$ )?
- ▶ **Computational Learning:** How efficient is the learning algorithm?

Typical learning frameworks:

- ▶ Estimation (e.g. coin tossing)
- ▶ Prediction (e.g. classification)
- ▶ Clustering
- ▶ Representation (Feature) Learning
- ▶ Density Estimation...

All the frameworks can be broadly generalized into two learning problems:

- ▶ Concept Learning (Binary Outcomes)
- ▶ Function Learning (Generalized Outcomes)

# Generalization 1: Concept Learning

- ▶ *Concept class*: A class  $\mathcal{C}$  of subsets of  $\mathcal{X}$
- ▶ *Unknown target concept*:  $C^* \in \mathcal{C}$  picked by Nature
- ▶ Binary Label:  $Y_i = \mathbb{1}_{\{X_i \in C^*\}}$
- ▶ The  $N$  feature-label pairs form the training set

$$\mathbf{Z}_1 = (X_1, Y_1) = \left(X_1, \mathbb{1}_{\{X_1 \in C^*\}}\right), \dots, \mathbf{Z}_N = (X_N, Y_N) = \left(X_N, \mathbb{1}_{\{X_N \in C^*\}}\right).$$

***The objective is to approximate target concept  $C^*$  as accurately as possible.***

Examples: Classification

## Problem 1: Concept Learning

- ▶ A concept learning problem is a triple  $(\mathcal{X}, \mathcal{P}, \mathcal{C})$ , where  $\mathcal{X}$  is the feature space,  $\mathcal{P}$  is a family of probability distributions on  $\mathcal{X}$ , and  $\mathcal{C}$  is a concept class.
- ▶ A learning algorithm for  $(\mathcal{X}, \mathcal{P}, \mathcal{C})$  is a sequence  $\mathcal{A} = \{A_n\}_{n=1}^\infty$  of mappings  $A_N : (\mathcal{X} \times \{0, 1\})^N \rightarrow \mathcal{C}$ .

Given a training set  $\mathbf{Z}_{1:N} = (Z_1, \dots, Z_N) \in \mathcal{Z}^N$  and a learning algorithm  $\mathcal{A}$ , the approximation to  $C^*$  is

$$\hat{C}_N = A_N(\mathbf{Z}_{1:N}) = A_N(Z_1, \dots, Z_N) = A_N\left(\left(X_1, \mathbb{1}_{\{X_1 \in C^*\}}\right), \dots, \left(X_N, \mathbb{1}_{\{X_N \in C^*\}}\right)\right).$$

# Generalization 1: Concept Learning (cont...)

Two types of errors: (i)  $X \in C^* \cap \hat{C}_N^c$ , (ii)  $X \in (C^*)^c \cap \hat{C}_N$ .

Combining the two, misclassification happens when  $X$  lies in the symmetric difference

$$C^* \Delta \hat{C}_N = (C^* \cap \hat{C}_N^c) \cup ((C^*)^c \cap \hat{C}_N).$$

Performance measure of  $\mathcal{A}$ :  $L(C^*, \hat{C}_N) = \mathbb{P}(C^* \Delta \hat{C}_N) = \mathbb{P}(X \in C^* \Delta \hat{C}_N)$ .

Good Algorithm  $\Rightarrow L(C^*, \hat{C}_N) \rightarrow 0$  as  $N \rightarrow \infty$ .

- ▶ Let  $X \sim P$ , and  $(X, \mathbb{1}_{X_N \in C}) \sim P_C$  for any  $C \in \mathcal{C}$ .
- ▶ Since  $\hat{C}_N$  is a random element in  $\mathcal{C}$ , the above convergence can only be achieved in a stochastic sense.
- ▶ Define “worst case” size of set of “bad” samples as

$$\phi_{\mathcal{A}}(N, \epsilon, P) = \sup_{C \in \mathcal{C}} P_C^N \left( L(C, A_N(\mathbf{Z}_{1:N})) > \epsilon \right)$$

- ▶ Since we do not know  $P$ , consider the worst case distribution as shown below:

$$\Phi_{\mathcal{A}}(N, \epsilon, \mathcal{P}) = \sup_{P \in \mathcal{P}} \phi_{\mathcal{A}}(N, \epsilon, P).$$

# Generalization 1: Concept Learning (cont...)

## Definition 1: PAC for Concept Learning

A learning algorithm  $\mathcal{A} = \{A_N\}$  is probably approximately correct<sup>a</sup> (or PAC) to accuracy  $\epsilon$  if

$$\lim_{N \rightarrow \infty} \Phi_{\mathcal{A}}(N, \epsilon, \mathcal{P}) = 0.$$

- ▶ We say that  $\mathcal{A}$  is PAC, if it is PAC to accuracy  $\epsilon$  for every  $\epsilon > 0$ .
- ▶ The concept class  $\mathcal{C}$  is called PAC-learnable to accuracy  $\epsilon$  w.r.t.  $P$ , if there exists an algorithm that is PAC to accuracy  $\epsilon$ .
- ▶ Finally, we say that  $\mathcal{C}$  is PAC-learnable, if there exists an algorithm that is PAC.

---

<sup>a</sup>D. Angluin. Queries and concept learning. Machine Learning, 2:319–342, 1988.

Equivalently, a learning algorithm  $\mathcal{A} = \{A_N\}$  is PAC, if for any  $\epsilon > 0$  and  $\delta > 0$ , there exists some  $N^*(\epsilon, \delta) \in \mathbb{N}$  such that, for all  $N \geq N^*(\epsilon, \delta)$ ,  $C \in \mathcal{C}$  and  $P \in \mathcal{P}$ , we have

$$P_C^N \left( L(C, A_N(\mathbf{Z}_{1:N})) > \epsilon \right) \leq \delta.$$

**Note:**  $N^*(\epsilon, \delta)$  is called the **sample complexity** of the learning algorithm  $\mathcal{A}$ .



# Example: Axis-Parallel Rectangles

- ▶ Let  $\mathcal{X} = [0, 1]^2$  and  $\mathcal{P}$  denote the set of all probability distributions on  $\mathcal{X}$
- ▶ Let  $\mathcal{C}$  denote the collection of all axis-parallel rectangles in  $\mathcal{X}$ , i.e.,  $C$  is in  $\mathcal{C}$  if it takes the form


$$C = [a_1, b_1] \times [a_2, b_2] = \left\{ (x_1, x_2) \in [0, 1]^2 \mid a_1 \leq x_1 \leq b_1, a_2 \leq x_2 \leq b_2 \right\},$$

for some  $0 \leq a_1 \leq b_1 \leq 1$  and  $0 \leq a_2 \leq b_2 \leq 1$ .

## Learning Algorithm:

Consider an intuitive algorithm  $\mathcal{A} = \{A_N\}_{N=1}^{\infty}$  where, for each  $N$ , we have

$$\begin{aligned}\hat{C}_N &= A_N(\mathbf{Z}_{1:N}) \\ &= \text{smallest rectangle } C \in \mathcal{C} \text{ that contains} \\ &\quad \text{all positive samples in } \mathbf{Z}_{1:N}.\end{aligned}$$



## Connections to Computer Vision:




This problem is tangentially similar to estimating bounding boxes in images. The number of samples in  $\mathcal{X}$  is similar to the resolution of an image.

# Example: Axis-Parallel Rectangles (cont...)

## Theorem 1

The above tightest rectangle algorithm  $\mathcal{A}$  is PAC, and therefore, the class  $\mathcal{C}$  is PAC-learnable, since

$$\Phi_{\mathcal{A}}(N, \epsilon, \mathcal{P}) \leq 4 \left(1 - \frac{\epsilon}{4}\right)^N.$$



**Proof of Theorem 1:**

## Example: Axis-Parallel Rectangles (cont...)

# Generalization 2: Function Learning

- ▶ *Function class*: A class  $\mathcal{F}$  defined on  $\mathcal{X}$
- ▶ *Target function*:  $f^* \in \mathcal{F}$  picked by nature
- ▶ Real-valued output:  $Y_i = f^*(X_i)$
- ▶ The  $N$  input-output pairs

$$Z_1 = (X_1, Y_1) = (X_1, f^*(X_1)), \dots, Z_N = (X_N, Y_N) = (X_N, f^*(X_N)).$$

***The objective is to approximate target function  $f^*$  as accurately as possible.***

Examples: Estimation

## Problem 2: Function Learning

- ▶ A function learning problem is a triple  $(\mathcal{X}, \mathcal{P}, \mathcal{F})$ , where  $\mathcal{X}$  is the feature space,  $\mathcal{P}$  is a family of probability distributions on  $\mathcal{X}$ , and  $\mathcal{F}$  is a class of functions  $f : \mathcal{X} \rightarrow [0, 1]$ .
- ▶ A learning algorithm for  $(\mathcal{X}, \mathcal{P}, \mathcal{F})$  is a sequence  $\mathcal{A} = \{A_n\}_{n=1}^\infty$  of mappings  $A_N : (\mathcal{X} \times \{0, 1\})^N \rightarrow \mathcal{F}$ .

Given a training set  $\mathbf{Z}_{1:N} = (Z_1, \dots, Z_N) \in \mathcal{Z}^N$  and a learning algorithm  $\mathcal{A}$ , the approximation of  $f^*$  is

$$\hat{f}_N = A_N(\mathbf{Z}_{1:N}) = A_N\left(\left(X_1, \mathbb{1}_{\{X_1 \in C^*\}}\right), \dots, \left(X_N, \mathbb{1}_{\{X_N \in C^*\}}\right)\right).$$

## Generalization 2: Function Learning (cont...)

Performance of  $\mathcal{A}$ :  $L_P(\hat{f}_N, f^*) = \mathbb{E}_P \left[ \left| \hat{f}_N - f^* \right|^2 \right] = \int_{\mathcal{X}} |\hat{f}_N(x) - f^*(x)|^2 P(dx)$

**Remark:** Concept learning is a special case of function learning.

- Define “worst case” size of set of “bad” samples as

$$\phi_{\mathcal{A}}(N, \epsilon, P) = \sup_{f \in \mathcal{F}} P_f^N \left( L(A_N(\mathbf{Z}_{1:N}), f) > \epsilon \right)$$

- Since we do not know  $P$ , consider the worst case distribution as shown below:

$$\Phi_{\mathcal{A}}(N, \epsilon, \mathcal{P}) = \sup_{P \in \mathcal{P}} \phi_{\mathcal{A}}(N, \epsilon, P).$$

### Definition 2: PAC for Function Learning

A learning algorithm  $\mathcal{A} = \{A_N\}$  is probably approximately correct (or PAC) to accuracy  $\epsilon$  if

$$\lim_{N \rightarrow \infty} \Phi_{\mathcal{A}}(N, \epsilon, \mathcal{P}) = 0.$$

- The function class  $\mathcal{F}$  is called PAC-learnable to accuracy  $\epsilon$  w.r.t.  $P$ , if there exists an algorithm that is PAC to accuracy  $\epsilon$ .
- Finally,  $\mathcal{F}$  is called PAC-learnable, if there exists an algorithm that is PAC.

# Limitations of Model-Based Approaches

- ▶ We assume  $C^* \in \mathcal{C}$  (or equivalently,  $f^* \in \mathcal{F}$ )  $\Rightarrow$  Fit data regarding a well-studied phenomenon to some **a priori known hypothesis class**
- ▶ Labels  $y = \mathbf{1}_{x \in C^*}$  (or equivalently,  $y = f^*(x)$ ) are assumed to be **noiseless**.

Such limitations will lead us naturally towards a new framework called **model-agnostic learning** (also called model-free learning).

***The main goal is to find the best possible hypothesis (concept/function) within a chosen hypothesis class  $\mathcal{F}$ .***

## Problem 3: Model-Agnostic Learning

A model-agnostic learning problem is a tuple  $(\mathcal{X}, \mathcal{Y}, \mathcal{P}, \mathcal{F})$ , where

- ▶ Sets:  $\mathcal{X}$  (input feature space),  $\mathcal{Y}$  (label space) and  $\mathcal{U}^a$  (hypothesis space)
- ▶ A class  $\mathcal{P}$  of probability distributions on  $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ .
- ▶ A class  $\mathcal{F}$  of functions  $f : \mathcal{X} \rightarrow \mathcal{U}$ .

A learning algorithm for  $(\mathcal{X}, \mathcal{Y}, \mathcal{U}, \mathcal{P}, \mathcal{F}, \ell)$  is a sequence of mappings  $\mathcal{A} = \{A_N\}_{N=1}^{\infty}$ , where  $A_N : \mathcal{Z}^N \rightarrow \mathcal{F}$ .

---

<sup>a</sup> $\mathcal{U} \neq \mathcal{Y}$ , since the true-hypothesis labels are corrupted by noise.

# PAC Learnability for Model-Agnostic Learning

- ▶ Given a learning algorithm  $\mathcal{A} = \{A_N\}_{N=1}^{\infty}$  with  $A_N : \mathcal{Z}^N \rightarrow \mathcal{F}$ , if  $\hat{f}_N = A_N(\mathbf{Z}_{1:N})$ , then the performance can be measured as

$$L_P(\hat{f}_N) = \mathbb{E}_P \left[ \ell(Y, \hat{f}_N(X)) \right] = \int_{\mathcal{X} \times \mathcal{Y}} \ell(y, \hat{f}_N(x)) P(dx, dy)$$

- ▶ **Minimum risk:**  $L_P^*(\mathcal{F}) = \inf_{f \in \mathcal{F}} L_P(f)$  for an induced function class  $\mathcal{F}$ .  
i.e., given any algorithm  $\mathcal{A}$ , we have  $0 \leq L_P^*(\mathcal{F}) \leq L_P(\hat{f}_N) \leq 1$ .
- ▶ Given any  $\epsilon > 0$ , let the worst case probability of getting a bad sample be

$$\Phi_{\mathcal{A}}(N, \epsilon) = \sup_{P \in \mathcal{P}} P^N \left( L_P(\hat{f}_N) > L_P^*(\mathcal{F}) + \epsilon \right)$$

## Definition 3: PAC for Model-Agnostic Learning

A learning algorithm  $\mathcal{A} = \{A_N\}$  for a problem  $(\mathcal{X}, \mathcal{Y}, \mathcal{U}, \mathcal{P}, \mathcal{F}, \ell)$  is PAC to accuracy  $\epsilon$  if

$$\lim_{N \rightarrow \infty} \Phi_{\mathcal{A}}(N, \epsilon) = 0.$$

- ▶ An algorithm that is PAC to accuracy  $\epsilon$  for every  $\epsilon > 0$  is said to be PAC.
- ▶ Finally, a learning problem  $(\mathcal{X}, \mathcal{Y}, \mathcal{U}, \mathcal{P}, \mathcal{F}, \ell)$  is *model-agnostically learnable* if there exists an algorithm for it, which is PAC.

# Empirical Risk Minimization

But, we do not always know the input distribution  $P \in \mathcal{P}$ .

- ▶  $L_P(f)$  is unknown. Can we replace this with some surrogate?
- ▶ **ERM Algorithm:**  $\hat{f}_N = \arg \min_{f \in \mathcal{F}} L_N(f) \triangleq \frac{1}{N} \sum_{i=1}^N \ell(Y_i, f(X_i))$

## Theorem 2: Fundamental Theorem of Learning

Consider an agnostic learning problem  $(\mathcal{X}, \mathcal{P}, \mathcal{F})$  and let  $\delta > 0$ . For any  $P \in \mathcal{P}$ , the ERM algorithm satisfies

$$L_P(\hat{f}_N) \leq L_P^*(\mathcal{F}) + 8\sqrt{\frac{V(\mathcal{F}) \log(n+1)}{n}} + \sqrt{\frac{2 \log(1/\delta)}{n}}$$

with probability at least  $1 - \delta$ .

Moreover, there is a universal constant  $C$  so that for any distribution  $P$  on  $\mathcal{Z}$  and  $\delta \in (0, 1)$ , the ERM algorithm satisfies

$$L_P(\hat{f}_N) \leq L_N(\hat{f}_N) + C\sqrt{\frac{V(\mathcal{F})}{n}} + \sqrt{\frac{2 \log(1/\delta)}{n}}$$

with probability at least  $1 - \delta$ .



# Vapnik-Chervonenkis Dimension

## Definition 4: Shattering

Let  $\mathcal{C}$  be a class of (measurable) subsets of some space  $\mathcal{Z}$ . We say that a finite set  $S = \{z_1, \dots, z_N\} \subset \mathcal{Z}$  is shattered by  $\mathcal{C}$ , if for every subset  $S' \subseteq S$ , there exists some  $C \in \mathcal{C}$  such that  $S' = S \cap C$ .

i.e.  $S = \{z_1, \dots, z_N\} \subset \mathcal{Z}$  is shattered by  $\mathcal{C}$  if, for any binary  $N$ -tuple  $b = (b_1, \dots, b_N) \in \{0, 1\}^N$ , there exists some  $C \in \mathcal{C}$  such that

$$(\mathbb{1}_{\{z_1 \in C\}}, \dots, \mathbb{1}_{\{z_N \in C\}}) = b.$$

$$N^{\text{th}} \text{ Shatter Coefficient of } \mathcal{C}: \quad \mathbb{S}_N(\mathcal{C}) = \sup_{S \subset \mathcal{Z}; |S|=n} |\{S \cap C \mid C \in \mathcal{C}\}|.$$

## Definition 5: VC Dimension

The Vapnik-Chervonenkis dimension (or the VC dimension) of  $\mathcal{C}$  is

$$V(\mathcal{C}) = \sup \{|S|, \text{ such that } S \text{ is shattered by } \mathcal{C}\}.$$

If  $V(\mathcal{C}) < \infty$ , we say that  $\mathcal{C}$  is a VC class (of sets).

In other words,

$$V(\mathcal{C}) = \sup \{n \in \mathbb{N} \mid \mathbb{S}_n(\mathcal{C}) = 2^n\}.$$

# Vapnik-Chervonenkis Dimension (cont...)

## Definition 6: VC Dimension of $\mathcal{F}$

Let  $\mathcal{F}$  be a class of functions  $f : \mathcal{Z} \rightarrow \{0, 1\}$ . We say that a finite set  $S = \{z_1, \dots, z_n\} \subset \mathcal{Z}$  is shattered by  $\mathcal{F}$  if it is shattered by the class  $\mathcal{C}_{\mathcal{F}} = \{C_f : f \in \mathcal{F}\}$ , where  $C_f = \{z \in \mathcal{Z} \mid f(z) = 1\}$ .

### Example 1: Semi-Infinite Intervals

Let  $\mathcal{Z} = \mathbb{R}$  and  $\mathcal{C}$  is a class of all intervals of the form  $(\infty, t)$ . Then,  $V(\mathcal{C}) = 1$ .

### Example 2: Closed Intervals

Let  $\mathcal{Z} = \mathbb{R}$  and  $\mathcal{C}$  is a class of all intervals of the form  $(s, t)$ . Then,  $V(\mathcal{C}) = 2$ .

### Example 3: Closed Half-Spaces

Let  $\mathcal{Z} = \mathbb{R}^2$  and  $\mathcal{C}$  is a class of all closed half-spaces, i.e. the sets of the form  $\{z = (z_1, z_2) \in \mathbb{R}^2 \mid w_1 z_1 + w_2 z_2 \geq b\}$  for all choices of  $w_1, w_2, b$  such that  $(w_1, w_2) \neq (0, 0)$ . Then,  $V(\mathcal{C}) = 3$ .

### Example 4: Axis-Parallel Rectangles

Let  $\mathcal{Z} = \mathbb{R}^2$  and  $\mathcal{C}$  is a class of all axis-parallel rectangles, i.e. the sets of the form  $C = [a_1, b_1] \times [a_2, b_2]$  for all  $a_1, a_2, b_1, b_2 \in \mathbb{R}$ . Then,  $V(\mathcal{C}) = 4$ .

# No-Free-Lunch Theorem

$$\mathbb{P} \left[ L_P(\hat{f}_N) \leq L_N(\hat{f}_N) + C \sqrt{\frac{V(\mathcal{F})}{n}} + \sqrt{\frac{2 \log(1/\delta)}{n}} \right] \geq 1 - \delta.$$

The converse of Theorem 2 is also true, i.e. *there is no universal learner*.

## Theorem 3: No-Free-Lunch Theorem

Let  $\mathcal{A} = \{A_N\}_{N \in \mathbb{N}}$  be any binary classification algorithm with 0–1 binary loss on a domain  $\mathcal{X}$ . Let the training set be of size  $m \leq \frac{|\mathcal{X}|}{2}$ . Then, there always exists a distribution  $P$  over  $\mathcal{X} \times \{0, 1\}$  such that

1. there exists a function  $f : \mathcal{X} \rightarrow \{0, 1\}$  such that  $L_P(f) = 0$ ,
2. with probability at least  $1/7$  over some training set  $S \sim P^m$ , we have  $L_P(A_m(S)) \geq 1/8$ .

**Interpretation:** Design ML algorithm that performs well on a specific task on a practical data distribution space.




# Perceptron<sup>2</sup>: A Linear Threshold Unit (LTU)

$$y = \sigma(\mathbf{w}^T \mathbf{x}) = \sigma\left(\sum_{i=0}^d w_i x_i\right) \in \{-1, +1\},$$

- ▶  $x_0 = 1$  (bias neuron)
- ▶  $w_0$  is the threshold
- ▶  $\sigma$  is the *activation/squashing* function.
- ▶ **Example:** Heavyside step function

$$\sigma(x) = \begin{cases} 1, & \text{if } x \geq 0, \\ 0, & \text{otherwise.} \end{cases}$$


- ▶ Mimics synapses in biological neural networks




<sup>2</sup>Frank Rosenblatt, "The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain," *Psychological Review*, vol. 65, no. 6, pp: 386, 1958.

# Perceptrons and Boolean Functions

## Can perceptrons model boolean functions?



Perceptrons can model several generalized boolean expressions (e.g. majority rule).




### Lemma 1: Minsky and Papert, 1969

A perceptron cannot model a XOR gate.

In other words, there is no combination of weights and threshold in the perceptron model such that the output resembles an XOR.

# Need for Multilayer Perceptrons (or, Neural Networks)

- ▶ Stack perceptrons to model complex functions
- ▶ A simple two-layer neural network (one hidden layer, one output layer) can model XOR function.



Let  $\mathcal{F}_L$  denote the set of all neural networks with  $L$  layers.

## Theorem 4

$\mathcal{F}_2$  is a universal Boolean function space.


### Caveat:

- ▶ An arbitrary Boolean function can be represented as a truth table.
- ▶ A truth table can be represented in disjunctive normal form (DNF)
- ▶ DNF expressions may need a very large number of neurons in the hidden layer.
- ▶ Reduce the number of neurons with the help of Karnaugh maps (K-Maps)


**Can we reduce the number of neurons?**

# Need for Deep Neural Networks (DNNs)

**Example<sup>3</sup>:** Reduction from 7 hidden neurons to 3 hidden neurons.



$$\begin{aligned} &\bar{W}\bar{X}\bar{Y}\bar{Z} + \bar{W}X\bar{Y}\bar{Z} + W\bar{X}\bar{Y}\bar{Z} \\ &+ W\bar{X}\bar{Y}Z + \bar{W}X\bar{Y}Z + \bar{W}\bar{X}Y\bar{Z} \\ &+ W\bar{X}Y\bar{Z} \end{aligned}$$




$$\bar{Y}\bar{Z} + \bar{W}X\bar{Y} + \bar{X}Y\bar{Z}$$


## Largest Irreducible DNF:

### Lemma 2

A perceptron may require  $O(2^N)$  hidden neurons to represent a Boolean function of  $N$  variables.



4 variables  $\Rightarrow$  8 terms







6 variables  $\Rightarrow$  32 terms

<sup>3</sup>This example was borrowed from the lecture notes of "Deep Learning" course offered in Spring 2022 by Bhiksha Raj and Rita Singh at Carnegie Mellon University

# Need for DNNs (cont...)

- ▶ XOR of two variables needs 3 neurons
- ▶ XOR of  $N$  variables  $\Rightarrow$   $(N - 1)$  XORs of two variables
- ▶  $N$  variables needs  $3(N - 1)$  neurons



## Lemma 3

A perceptron requires  $O(N)$  hidden neurons to represent an XOR function of  $N$  variables.

- ▶ Neurons can be arranged in  $2 \log_2(N)$  layers.

## Claim 2

Reducing the number of layers below the minimum will result in an exponentially sized network for full representation.



# An Important Note...


## Claim 3: Shannon's Theorem

For  $N > 2$ , there is a Boolean function of  $N$  variables that requires at least  $\frac{2^N}{N}$  Boolean gates.



Note: If we could represent every Boolean function in  $O(N)$  hidden neurons, then  $P = NP$ !

# NNs as Universal Classifiers

- Each boundary line is a neuron.
- Convex region  $\Rightarrow$  AND of all boundary neurons
- Union of all convex regions  $\Rightarrow$  OR gate.




Circular regions can be approximated with higher-order polygons – *Polygon Nets*



# NNs as Universal Classifiers (cont...)

A good approximation of a circle needs asymptotic number of neurons!



Any arbitrary region is a union/intersection of circular regions (topology on  $\mathcal{X}$ ).

## Theorem 5


Neural networks are universal classifiers

**Note:** Deeper networks require far fewer neurons in most classifiers.



# Architecture Design for Complete Representation<sup>4</sup>

When is a NN architecture sufficiently broad/deep to represent a function?

- ▶ The adjacent pattern is composed of 16 lines
- ▶ So, a network with 16 or more neurons in the first layer is sufficient to represent the region.
- ▶ Also, need 40 neurons in the second layer.





- ▶ What if, we only have 8 neurons in the first layer?





<sup>4</sup>This example was borrowed from the lecture notes of "Deep Learning" course offered in Spring 2022 by Bhiksha Raj and Rita Singh at Carnegie Mellon University

# Architecture Design for Complete Representation (cont...)



- What if, we consider a different combination of lines?











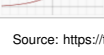
- **Note:** That information, which is not captured by existing neurons, is lost forever in all subsequent layers in the NN.
- This is mainly because of the heavyside step function!



- What if, we have a different activation function?



# Activation Functions to our Rescue!

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

Source: <https://towardsdatascience.com/activation-functions-neural-networks-1cbd9f8d91d6>

# More about Activation Functions...

- ▶ **Rectified Linear Unit (ReLU):**  $f(x) = \begin{cases} x, & \text{if } x \geq 0, \\ 0, & \text{otherwise.} \end{cases}$

- ▶ Good for intermediate layers
- ▶ *Concern:* No information on one side of hyperplane

- ▶ **Leaky Rectified Linear Unit (Leaky ReLU):**  $f(x) = \begin{cases} x, & \text{if } x \geq 0, \\ \alpha x, & \text{otherwise,} \end{cases}$  for any  $0 < \alpha \ll 1$ .

- ▶ Ideal for intermediate layers

- ▶ **Sigmoid (or logistic) function:**  $f(x) = \frac{1}{1 + e^{-x}}$

- ▶ Suits to estimate probability of an outcome in binary classification settings

- ▶ **Softmax (or generalized logistic) function:**  $f_i(x) = \frac{e^{x_i}}{\sum_{j=1}^M e^{x_j}}$ , for all

$$i = 1, \dots, M$$

- ▶ Suits to estimate probability of an outcome in  $M$ -ary classification settings

- ▶ **Hyperbolic tangent function:**  $f(x) = \tanh x$

- ▶ Suits to predict the outcome in binary classification settings

# NNs for Universal Function Approximation

## Definition 7: Universal Approximator

A class of functions  $\mathcal{F}$  is a universal approximator over a compact set  $S$ , if for every continuous function  $g$  and target accuracy  $\epsilon > 0$ , there exists  $f \in \mathcal{F}$  with

$$\sup_{x \in S} |f(x) - g(x)| \leq \epsilon.$$

Let  $\mathcal{F}_{\sigma,d}$  denote the set of all multilayer feedforward NNs  $\mathcal{F}$ , that is restricted to a  $d$ -dimensional input and uses  $\sigma(\cdot)$  as an activation function in all layers.

## Theorem 6: Hornik, Stinchcombe and White, 1989

Suppose  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  denote a continuous sigmoidal function such that



$$\lim_{z \rightarrow -\infty} \sigma(z) = 0, \quad \text{and} \quad \lim_{z \rightarrow \infty} \sigma(z) = 1.$$

Then, for any  $d \in \mathbb{R}$ ,  $\mathcal{F}_{\sigma,d}$  is universal.



# Radial Basis Function Networks for Function Approximation

Why always compare a point to some hyperplane?



\* Source: S. Cuomo, A. Galletti, G. Giunta and A. Starace, "Surface reconstruction from scattered point via RBF interpolation on GPU," 2013 Federated Conference on Computer Science and Information Systems, pp. 433-440, 2013.

# Feedforward NN: A Formal Model

## Definition 8: Feedforward NN

A feedforward neural network is a directed acyclic graph  $\mathcal{G} = (V, E)$ , and a weight function  $w : E \rightarrow \mathbb{R}$ . Each node is a single neuron (LTU) which is modeled by an activation function  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ .

Let the set of nodes be decomposed into a union of disjoint nodes, i.e.

$$V = \bigcup_{\ell=0}^L V_{\ell},$$

such that each edge  $e \in E$  connects some node in  $V_{\ell-1}$  to  $V_{\ell}$  for some  $\ell = 1, \dots, L$ .

Let  $\mathbf{f}_{\ell}(\mathbf{x})$  denote the output of  $V_{\ell}$ , when the NN is fed with some input  $\mathbf{x}$ . Then,

$$\mathbf{f}_{\ell}(\mathbf{x}) = \sigma_{\ell} \left( W_{\ell} \cdot \mathbf{f}_{\ell-1}(\mathbf{x}) \right)$$

In other words,

$$\mathcal{F}_L = \left\{ \mathbf{x} \rightarrow \sigma_L \left( W_L \cdot \sigma_{L-1} \left( \dots \sigma_1 (W_1 \mathbf{x}) \dots \right) \right) \mid \|W_{\ell}\| \leq B \text{ for all } \ell = 1, \dots, L \right\}$$

# VC Dimension for Neural Networks

## Theorem 7

VC dimension of perceptrons ( $\mathcal{F}_1$  with a heavyside-step activation function) on  $\mathbb{R}^d$  is  $d + 1$ .

## Theorem 8: (Bartlett, Harvey, Liaw, Mehrabian 2019)

Let  $\sigma$  be a piecewise linear function,  $W$  be the number of weight parameters,  $L$  be the number of layers. Then,

$$V(\mathcal{F}_L) \leq O(WL \log W).$$

Furthermore, there also exist networks with  $V(\mathcal{F}_L) \geq \Omega(WL \log(W/L))$ .

# Training Neural Networks: ERM Algorithm

- ▶ Training data:  $z_1 = (\mathbf{x}_1, y_1), \dots, z_N = (\mathbf{x}_N, y_N)$
- ▶ Function space: Set of all  $L$ -layered neural networks  $\mathcal{F}_L$ 
  - ▶ Activation functions  $\sigma_1(\cdot), \dots, \sigma_L(\cdot)$  are fixed.
  - ▶  $M$  outputs in one-hot form, for  $M$ -ary classification.
- ▶ Empirical Loss function:  $L_N(\mathbb{W}|\mathbf{z}_{1:N}) = \frac{1}{N} \sum_{i=1}^N \ell(y_i, f(\mathbf{x}_i|\mathbb{W}))$ 
  - ▶  $L_p$ -norm:  $\ell(y, f(\mathbf{x}|\mathbb{W})) = \left( \sum_{m=1}^M (y_m - f_m(\mathbf{x}|\mathbb{W}))^p \right)^{\frac{1}{p}}$
  - ▶ Binary Cross Entropy:  $\ell(y, f(\mathbf{x}|\mathbb{W})) = -y \log f(\mathbf{x}|\mathbb{W}) - (1 - y) \log (1 - f(\mathbf{x}|\mathbb{W}))$
  - ▶ Cross Entropy (or KL Divergence):  $\ell(y, f(\mathbf{x}|\mathbb{W})) = \sum_{m=1}^M y_m \log f_m(\mathbf{x}|\mathbb{W})$  for  $M$  classes, where  $y$  and  $f(\cdot)$  are represented in one-hot form.

**ERM Algorithm:** minimize  $L_N(\mathbb{W}|\mathbf{z}_{1:N})$   
 $\mathbb{W}$

**How do we solve the above optimization problem?**

# Gradient Descent




- Change  $\mathbb{W}$  in the opposite direction of gradient of  $L_N$ .

$$\mathbb{W}^{(r)} = \mathbb{W}^{(r-1)} - \delta \cdot \text{sgn}\left(\nabla L_N\left(\mathbb{W}^{(r-1)}\right)\right)$$

GD-SIGN( $\nabla L_N, \mathbb{W}^{(0)}, R$ )

```
1 for  $r = 1$  to  $R$   
2      $\mathbb{W}^{(r)} = \mathbb{W}^{(r-1)} - \delta \cdot \text{sgn}\left(\nabla L_N\left(\mathbb{W}^{(r-1)}\right)\right)$   
3 return  $\mathbb{W}^{(R)}$ 
```

**Toy Example:** Gradient descent with fixed step size on  $f(x) = x^2$  looks like...






# Gradient Descent (cont...)

- Faster convergence  $\Rightarrow$  Adapt step size according to the gradient.

$\text{GD}(\nabla L_N, \mathbb{W}^{(0)})$

```
1  for  $r = 1$  to  $R$   
2       $\mathbb{W}^{(r)} = \mathbb{W}^{(r-1)} - \delta \cdot \nabla L_N(\mathbb{W}^{(r-1)})$   
3  return  $\mathbb{W}^{(R)}$ 
```

- **Toy Example:** Gradient descent algorithm on  $f(x) = x^2$  looks like...



# Convergence of Gradient Descent

## Theorem 9

Suppose  $L_N : \mathbb{R}^W \rightarrow \mathbb{R}$  is convex and Lipschitz continuous with constant  $\eta > 0$ , i.e. we have  $\|\nabla L_N(\mathbb{W}_1) - \nabla L_N(\mathbb{W}_2)\|^2 \leq \eta \|\mathbb{W}_1 - \mathbb{W}_2\|^2$  for any  $\mathbb{W}_1, \mathbb{W}_2 \in \mathbb{R}^W$ .

Then, the gradient descent algorithm after  $r$  rounds with a fixed step size  $\delta \leq 1/\eta$  will yield a solution  $\mathbb{W}^{(r)}$  which satisfies

$$L_N(\mathbb{W}^{(r)}) - L_N(\mathbb{W}^*) \leq \frac{\|\mathbb{W}^{(0)} - \mathbb{W}^*\|^2}{2r\delta},$$

where  $\mathbb{W}^* = \arg \min_{\mathbb{W} \in \mathcal{F}_L} L_N(\mathbb{W})$  is the optimal solution to ERM.

**Proof:**

# Gradient Descent: Variants and Limitations

- Rather than iteratively repeat the update step  $R$  times, define a convergence criteria based on loss function...

GD2( $\nabla L_N, \mathbb{W}^{(0)}, \epsilon$ )

```
1   $r = 1$ 
2  while  $|L_N(\mathbb{W}^r) - L_N(\mathbb{W}^{r-1})| > \epsilon$ 
3       $\mathbb{W}^{(r+1)} = \mathbb{W}^{(r)} - \delta \cdot \nabla L_N(\mathbb{W}^{(r)})$ 
4       $r = r + 1$ 
5  return  $\mathbb{W}^{(r)}$ 
```

**Can we achieve faster convergence?**

- Note that most practical neural network architectures (wide or deep) have a large number of weight parameters.
- Computing the gradient is a challenging task.
- What if, we only compute the gradient in just one random direction?




# Stochastic Gradient Descent

- Compute gradient only in one random direction (say  $j \in \{1, \dots, W\}$ )
- Let  $[\nabla_{\mathbb{W}} L_N(\mathbb{W}^{(r-1)})]_j = \left[ 0 \quad \dots \quad 0 \quad \frac{\partial L_N(\mathbb{W}^{(r-1)})}{\partial \mathbb{W}_j} \quad 0 \quad \dots \quad 0 \right]^T$ .
- Requires significantly more rounds to converge
- However, each round is extremely fast!

SGD( $\nabla L_N, \mathbb{W}^{(0)}, \epsilon$ )

- 1  $r = 1$
- 2 **while**  $|L_N(\mathbb{W}^r) - L_N(\mathbb{W}^{r-1})| > \epsilon$
- 3     Pick some  $j \in \{1, \dots, W\}$  at random
- 4      $\mathbb{W}^{(r)} = \mathbb{W}^{(r-1)} - \delta \cdot [\nabla_{\mathbb{W}} L_N(\mathbb{W}^{(r-1)})]_j$
- 5      $r = r + 1$
- 6 **return**  $\mathbb{W}^{(r)}$



# Accelerated Gradient Descent and Nesterov's Momentum<sup>5</sup>

- ▶ GD traverses very slowly on a nearly flat surface.
- ▶ **Note:** Mimic how a ball gains its momentum as it rolls down a nearly horizontal incline.

- ▶ **Gradient Descent with Momentum (GDM):**


$$\mathbb{W}^{(r+1)} = \mathbb{W}^{(r)} + \mathbf{m}^{(r+1)}$$

$$\mathbf{m}^{(r+1)} = \beta \cdot \mathbf{m}^{(r)} - \alpha \cdot \nabla L_N(\mathbb{W}^{(r)})$$

- ▶ If  $\beta = 0$ , we have GD.
- ▶ **Problem:** GDM does not slow down enough at the bottom of the valley, and tends to overshoot the floor value.
- ▶ **Solution:** Gradient Descent with Nesterov's Momentum (GDNM)
  - ▶ Use the gradient at the projected future position instead...

$$\mathbb{W}^{(r+1)} = \mathbb{W}^{(r)} + \mathbf{m}^{(r+1)}$$

$$\mathbf{m}^{(r+1)} = \beta \cdot \mathbf{m}^{(r)} - \alpha \cdot \nabla L_N(\mathbb{W}^{(r)} + \beta \cdot \mathbf{m}^{(r)})$$



Source: M. J. Kochenderfer and T. A. Wheeler,  
"Algorithms for Optimization," The MIT Press, 2019.

<sup>5</sup>Y. Nesterov, "A Method of Solving a Convex Programming Problem with Convergence Rate  $O(1/k^2)$ ," *Soviet Mathematics Doklady*, vol. 27, no. 2, pp. 543-547, 1983.

# AdaGrad<sup>6</sup> and RMSProp

Can we update different components of  $\mathbf{W}$  using a different learning rate?



---

<sup>6</sup>J. Duchi, E. Hazan, and Y. Singer, "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization," *JMLR*, vol. 12, pp. 2121-2159, 2011.