

## **Topic 2: Efficient Learning**

# Inefficiencies in Learning

- ▶ Sampling Error and Overfitting
  - ▶ What if, the training data is sampled to have a certain pattern that is not present in the original input-output relationship?
- ▶ Variance in Stochastic Gradient Descent
  - ▶ Randomness in Gradient estimates can introduce a bias that is directly proportional to the variance in the gradient estimate.
- ▶ Complexity of Gradient Computation
  - ▶ Gradients for each complex model is very tedious to compute.

# Can We Mitigate Overfitting?

- ▶ **Best Approach:** Get more data!
  - ▶ Train on different *bags* of data
- ▶ Use the right model...
  - ▶ Hard to accomplish... similar to model-based learning.
- ▶ Consider **model ensembles**...
  - ▶ Use different function classes (models of different forms)
  - ▶ Identify multiple weight parameters for different initializations and take their average
- ▶ How about **regularization** to limit the capacity of neural networks?
  - ▶ Limit the number of hidden layers and/or units per layer.
  - ▶ Early stopping criteria in optim algorithms, before overfitting begins
  - ▶ Penalize the objective for large weights using  $\ell_1$  penalty, or  $\ell_2$  penalty
  - ▶ Introduce hard constraint on weight capacity (*a.k.a.* max-norm, i.e.  $\|\mathbb{W}\|_p \leq c.$ ).

# Bootstrap Aggregation (in short, Bagging)<sup>1</sup>

- ▶ Reduce generalized error by aggregating the outcomes of several models.
- ▶ Generate a bootstrap sample, say  $\mathcal{D}_k$ , which follows the same distribution as the training set  $\mathcal{D}$ , for  $k = 1, \dots, K$ .
- ▶ Train a new classifier, say  $\hat{y}_k = \hat{f}_k(x)$ , on each bootstrap sample  $\mathcal{D}_k$ .
- ▶ **Aggregation rule:** Count the number of times a class label appears amongst the  $K$  classifiers. The label with highest count is returned as output. Ties are broken by choosing the labels with lowest class label.
- ▶ Can also aggregate via averaging the outcomes.
- ▶ Let  $\epsilon_k$  denote the error in the  $k^{th}$  classifier, with variance  $\mathbb{E}(\epsilon_k^2) = v$ , and covariance  $\mathbb{E}(\epsilon_k \epsilon_j) = c$ .

$$\begin{aligned}\mathbb{E} \left[ \left( \frac{1}{K} \sum_{k=1}^K \epsilon_k \right)^2 \right] &= \frac{1}{K^2} \mathbb{E} \left[ \sum_{k=1}^K \left( \epsilon_k^2 + \sum_{j \neq k} \epsilon_k \epsilon_j \right) \right] \\ &= \frac{1}{K} v + \frac{K-1}{K} c\end{aligned}$$


- ▶ If  $c = 0$ , squared error reduces by a factor of  $K$ .

---


<sup>1</sup>Leo Breiman, "Bagging Predictors," *Machine learning*, vol. 24, no. 2, pp. 123-140, 1996.

# Dropout<sup>2</sup>


- ▶ Bagging and other ensemble methods is hard, especially with large neural network models.
- ▶ **Simple Approach:** Drop neurons with a fixed probability  $p$ , during training.
- ▶ Aggregate by averaging the weight parameters across different models




(a) Standard Neural Net



(b) After applying dropout.



(a) At training time




(b) At test time

<sup>2</sup>N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, "Dropout: a simple way to prevent neural networks from overfitting," *The Journal of Machine Learning Research*, vol. 15, no. 1, pp. 1929-1958, 2014.

# Dropout (cont...)


Formally, dropout is modeled as follows...

$$\begin{aligned} z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \mathbf{y}^l + b_i^{(l+1)}, \\ y_i^{(l+1)} &= f(z_i^{(l+1)}), \end{aligned}$$



(a) Standard network

$$\begin{aligned} r_j^{(l)} &\sim \text{Bernoulli}(p), \\ \tilde{\mathbf{y}}^{(l)} &= \mathbf{r}^{(l)} * \mathbf{y}^{(l)}, \\ z_i^{(l+1)} &= \mathbf{w}_i^{(l+1)} \tilde{\mathbf{y}}^l + b_i^{(l+1)}, \\ y_i^{(l+1)} &= f(z_i^{(l+1)}). \end{aligned}$$



(b) Dropout network

- **Note:** Backpropagation on thinned networks over several mini-batches of data
- Dropout + Norm-Regularization + AdaM (with large learning rates and high momentum)  $\Rightarrow$  Huge boost in performance!

# Regularization in Neural Networks

$$\hat{f}_N = \arg \min_{\mathbb{W}} L_N(\mathbb{W}) + \lambda \Phi(\mathbb{W}),$$

where  $\Phi(\mathbb{W})$  is the term that penalizes undesirable weights.

- ▶ **Weight Decay ( $\ell_2$ -Penalty):**  $\Phi(\mathbb{W}) = \frac{1}{2} ||\mathbb{W}||_2^2$ 
  - ▶ Also called ridge regression, or Tikhonov regularization
  - ▶ Drives weights closer to zero.
  - ▶ Let  $\Psi_2(\mathbb{W}) = L_N(\mathbb{W}) + \lambda ||\mathbb{W}||_2^2 \Rightarrow \nabla \Psi_2(\mathbb{W}) = \nabla L_N(\mathbb{W}) + 2\lambda \mathbb{W}$
  - ▶ GD:  $\mathbb{W}^{r+1} = (1 - 2\delta\lambda)\mathbb{W}^r - \delta \nabla L_N(\mathbb{W}) \Rightarrow \mathbb{W} \rightarrow -\frac{1}{2\lambda} \nabla L_N(\mathbb{W})$
- ▶  **$\ell_1$ -Penalty:**  $\Phi_1(\mathbb{W}) = ||\mathbb{W}||_1$ 
  - ▶ Results in sparse  $\mathbb{W}$ .
  - ▶ Let  $\Psi_1(\mathbb{W}) = L_N(\mathbb{W}) + \lambda ||\mathbb{W}||_1 \Rightarrow \nabla \Psi_1(\mathbb{W}) = \nabla L_N(\mathbb{W}) + 2\lambda \cdot \text{sign}(\mathbb{W})$
  - ▶ Weight update:  $\mathbb{W}^{r+1} = \mathbb{W}^r - \delta \nabla L_N(\mathbb{W}) - 2\lambda\delta \cdot \text{sign}(\mathbb{W})$  (for each entry in  $\mathbb{W}$ )
    - ▶  $\mathbb{W}^{r+0.5} = \mathbb{W}^r - \delta \nabla L_N(\mathbb{W})$
    - ▶ If  $\mathbb{W}^{r+0.5} > 0$ , then  $\mathbb{W}^{r+1} = \max\{0, \mathbb{W}^{r+0.5} - 2\lambda\delta\}$
    - ▶ If  $\mathbb{W}^{r+0.5} < 0$ , then  $\mathbb{W}^{r+1} = \min\{0, \mathbb{W}^{r+0.5} + 2\lambda\delta\}$

# Batch Normalization<sup>3,4</sup>

- ▶ **Internal Covariance Shift:** Change in distribution of activations due to change in model parameters – slows training
- ▶ **Solution:** Whitening activation functions
  - ▶ Linearly transform features to have zero mean and unit variances, i.e. decorrelated.

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma,\beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma,\beta}(x_i) \quad // \text{ scale and shift}$$

---

<sup>3</sup>S. Ioffe, and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," in *International Conference on Machine Learning*, pp. 448-456, 2015.

<sup>4</sup>Ping Luo, Xinjiang Wang, Wenqi Shao, Zhanglin Peng, "Towards Understanding Regularization in Batch Normalization," *ICLR*, 2019.



# Batch Normalization (cont...)

## Backpropagation for BatchNorm:

$$\frac{\partial \ell}{\partial \hat{x}_i} = \frac{\partial \ell}{\partial y_i} \cdot \gamma$$

$$\frac{\partial \ell}{\partial \sigma_{\mathcal{B}}^2} = \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot (x_i - \mu_{\mathcal{B}}) \cdot \frac{-1}{2} (\sigma_{\mathcal{B}}^2 + \epsilon)^{-3/2}$$

$$\frac{\partial \ell}{\partial \mu_{\mathcal{B}}} = \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

$$\frac{\partial \ell}{\partial x_i} = \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} + \frac{\partial \ell}{\partial \sigma_{\mathcal{B}}^2} \cdot \frac{2(x_i - \mu_{\mathcal{B}})}{m} + \frac{\partial \ell}{\partial \mu_{\mathcal{B}}} \cdot \frac{1}{m}$$

$$\frac{\partial \ell}{\partial \gamma} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i} \cdot \hat{x}_i$$

$$\frac{\partial \ell}{\partial \beta} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i}$$

# Mini-Batch Stochastic Gradient Descent

- ▶ Random coordinate descent<sup>5</sup> (RCD) provides same convergence rate as that of full gradient descent
- ▶ RCD takes a fraction of effort by evaluating gradients only for one (or a small subset of) random coordinate(s).
- ▶ But, computer vision problems require a lot of images for training – RCD is *not* sufficient!
- ▶ **Solution:** Mini-batch SGD<sup>6</sup>

$$[\nabla L_N(\mathbb{W}^{r-1})]_{\mathcal{B}} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \ell(y_i, \hat{y}_i | \mathbb{W}^{r-1})$$

- ▶ Estimate the gradient using a small sample (mini-batch) of training data.
- ▶ Introduces bias which is a function of variance in the gradient estimate.
  - ▶ Variance in gradient estimate reduces as the size of mini-batch increases.
  - ▶ Alternatively, change the learning rate to  $\delta_t = \delta \cdot \gamma^t$ , where  $\gamma$  is a discounting factor that forces mini-batch SGD to converge.


---

<sup>5</sup>Y. Nesterov, "Efficiency of Coordinate Descent Methods on Huge-Scale Optimization Problems," *SIAM Journal on Optimization*, vol. 22, no. 2, pp. 341-362, 2012.

<sup>6</sup>M. Li, T. Zhang, Y. Chen, and A. J. Smola, "Efficient Mini-Batch Training for Stochastic Optimization," in *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 661-670. 2014.

# Data Parallelism: Accelerating Computation using Hardware


**Parallelizing Mini-Batch SGD:** Compute each mini-batch on a different GPU.




- ▶ Minibatch SGD is robust to *asynchronous* gradient computation
- ▶ However, some GPUs (workers) may compute gradient for older models due to asynchrony
- ▶ **Solution:** Synchronous SGD (Sync-SGD)
  - ▶ Wait until all gradients are available
  - ▶ Update time dictated by slowest worker.

# Model Parallelism - for Large (Wide/Deep) Neural Networks

**DistBelief:** Each model partition on a different CPU.




**Model Aggregation:** Models are trained on each mini-batch on a different GPU.



Source: J. Dean, G. Corrado, R. Monga, K. Chen, M. Devin, M. Mao, M. Ranzato et al. "Large Scale Distributed Deep Networks," *Advances in neural information processing systems*, vol. 25, 2012.

# Synchronous vs. Asynchronous Computation

- ▶ Optimization algorithms are robust to asynchronous gradient computation
- ▶ However, model parallelism is sensitive to asynchrony
- ▶ Synchronous Computation
  - ▶ Accurate outcomes
  - ▶ Slow due to some *stragglers* in the tail end of queue (gradients for lower layers).

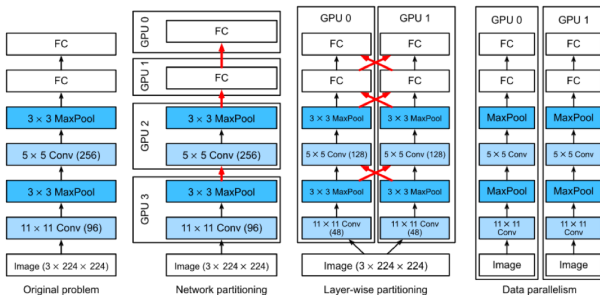


- ▶ **Fix:** Can be accelerated using backup GPUs and discarding straggler computations.
- ▶ Mixed Computation
  - ▶ Group certain number of workers together, synchronize intra group and perform asynchronous updates across groups<sup>7</sup>
  - ▶ Slower than Sync-SGD!


<sup>7</sup> J. Chen, R. Monga, S. Bengio, and R. Jozefowicz, "Revisiting Distributed Synchronous SGD," in *Proceedings of ICLR Workshop Track*, 2016. Available: <https://arxiv.org/abs/1604.00981>.

# Combining Model/Data Parallelisms...

In summary, there are several parallel architectures:



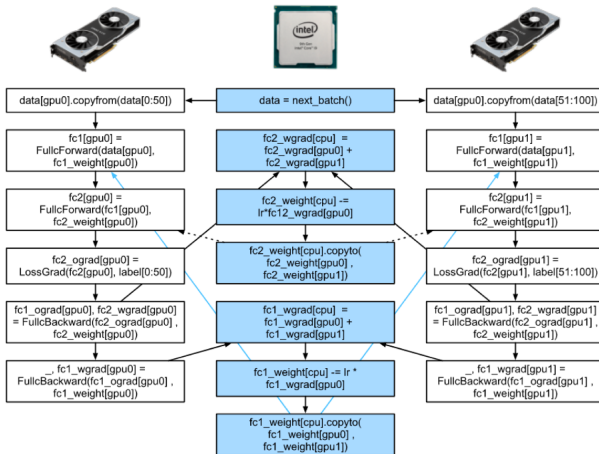
Accordingly, modern deep learning libraries use one of the communication strategies:



Source: M. Abadi *et al.*, "TensorFlow: A System for Large-Scale Machine Learning," in *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*, pp. 265-283, 2016.

# Illustrative Example: Data Workflow

Consider a simple two-layer neural network model trained on one CPU and two GPUs:



**The only difficult aspect is to compute gradients using backprop!**

# Computational Graphs and Automatic Differentiation<sup>8</sup>

- ▶ Automatic differentiation is **not** symbolic computation!
  - ▶ The goal is not to evaluate a formula – instead compute the gradients efficiently.
  - ▶ Identify a sequence of primitive computations, which have specific routines for computing derivatives
  - ▶ Perform backprop in a mechanical way.
- ▶ Types of Automatic differentiation
  - ▶ Forward Mode (slower method)
  - ▶ Reverse Mode – Backprop!
- ▶ Consider the following simple example of an NN-1 model for scalar input:

## Computation Sequence:

### Original Loss Function:

$$\ell = \frac{1}{2} \left[ \frac{1}{1 + \exp(-wx - b)} - y \right]^2$$

$$z = wx + b$$

$$\hat{y} = \frac{1}{1 + \exp(-z)}$$

$$\ell = \frac{1}{2} (\hat{y} - y)^2$$

- ▶ Automatic Differentiation  $\Rightarrow$  Construct a computational graph and evaluate backprop using primitives' gradients


---

<sup>8</sup>Source: A. G. Baydin, B. A. Pearlmutter, A. A. Radul, and J. M. Siskind. "Automatic Differentiation in Machine Learning: A Survey." *Journal of Machine Learning Research*, vol. 18, pp. 1-43, 2018.  
CS 6406: Machine Learning for Computer Vision (Sid Nadendla)



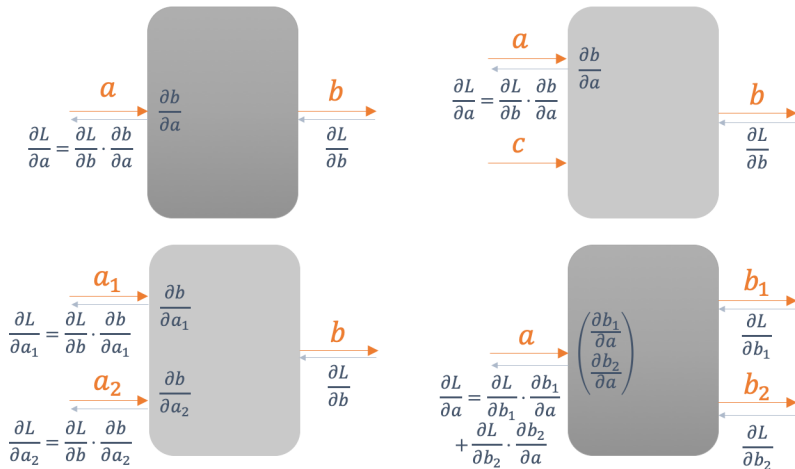
# Message Passing in Computational Graphs

- ▶ During the backward pass, each node in the graph receives upstream gradients
- ▶ Compute downstream gradients by multiplying them by local gradients



# Message Passing on Primitives

- Identify basic primitives that constitutes a given function
- Construct a computational graph automatically
- Perform message passing on gradients in reverse mode.



# How to Construct a Computational Graph?

Approach:

- ▶ Define a node class
- ▶ Construct a primitive
- ▶ Define primitive on basic numpy functions

```
add_new = primitive(np.add)
mul_new = primitive(np.multiply)
div_new = primitive(np.divide)
sub_new = primitive(np.subtract)
neg_new = primitive(np.negative)
exp_new = primitive(np.exp)
```

Ref. <https://tomroth.com.au/compgraph2/>

```
class Node:
    """A node in a computation graph."""
    def __init__(self, value, fun, parents):
        self.parents = parents
        self.value = value
        self.fun = fun

    def __repr__(self):
        """A (very) basic string representation"""
        if self.value is None: str_val = 'None'
        else: str_val = str(round(self.value,3))
        return "\n" + "Fun: " + str(self.fun) + \
            " Value: " + str_val + \
            " Parents: " + str(self.parents)

from functools import wraps
def primitive(f):
    @wraps(f)
    def inner(*args, **kwargs):
        ## Code to add operation/primitive to computation graph

        # We need to separate out the integer/non node case.
        Sometimes you are adding
        # constants to nodes.
        def getval(o): return o.value if type(o) == Node else o

        if len(args): argvals = [getval(o) for o in args]
        else: argvals = args
        if len(kwargs): kwargvals = dict([(k,getval(o)) for k,o in kwargs.items()])
        else: kwargvals = kwargs

        # get parents
        l = list(args) + list(kwargs.values())
        parents = [o for o in l if type(o) == Node ]

        value = f(*argvals, **kwargvals)
        print("add", "" + f.__name__ + "", "to graph with",value)
        return Node(value, f, parents)
    return inner
```

# Using these primitives...

```
def start_node(value = None):
    """A function to create an empty node to start off the
    graph"""
    fun,parents = lambda x: x, []
    return Node(value, fun, parents)

z = start_node(1.5)
t1 = mul_new(z, -1)
t2 = exp_new(t1)
t3 = add_new(t2, 1)
y = div_new(1,t3)
print("Final answer:", round(y.value,3)) # correct final output
print(y)
```

```
add 'multiply' to graph with value -1.5
add 'exp' to graph with value 0.22313016014842982
add 'add' to graph with value 1.22313016014843
add 'true_divide' to graph with value 0.8175744761936437
Final answer: 0.818

Fun: <ufunc 'true_divide'> Value: 0.818 Parents: [
Fun: <ufunc 'add'> Value: 1.223 Parents: [
Fun: <ufunc 'exp'> Value: 0.223 Parents: [
Fun: <ufunc 'multiply'> Value: -1.5 Parents: [
Fun: <function start_node.<locals>.<lambda> at 0x10fea27b8> Value:
1.5 Parents: []]]]]
```

# Computational Graphs in PyTorch: A Simple Example<sup>9</sup>

Consider a simple function:  $z = \log(x_1 x_2) * \sin(x_2)$

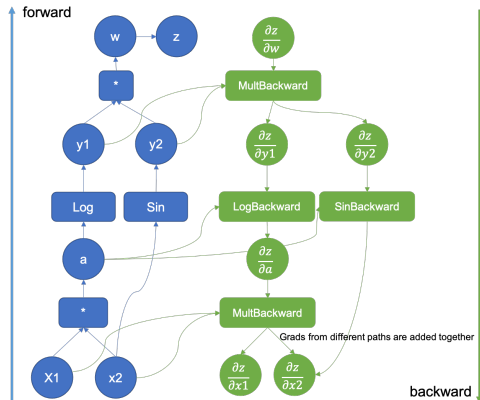
## Primitives:

- ▶  $a = x_1 x_2$
- ▶  $y_1 = \log(a)$
- ▶  $y_2 = \sin(x_2)$
- ▶  $w = y_1 y_2$
- ▶  $z = w$

```
>>> import torch
>>> x = torch.tensor([0.5, 0.75], requires_grad=True)
>>> y = torch.log(x[0] * x[1]) * torch.sin(x[1])
>>> y.backward(1.0)
>>> x.grad
tensor([1.3633, 0.1912])
```

## Derivatives:

- ▶  $\frac{\partial z}{\partial w} = 1$
- ▶  $\frac{\partial z}{\partial y_1} = \frac{\partial z}{\partial w} \cdot \frac{\partial w}{\partial y_1} = 1 \cdot y_2$
- ▶  $\frac{\partial z}{\partial y_2} = \frac{\partial z}{\partial w} \cdot \frac{\partial w}{\partial y_2} = 1 \cdot y_1$
- ▶  $\frac{\partial z}{\partial a} = \frac{\partial z}{\partial y_1} \cdot \frac{\partial y_1}{\partial a} = y_2 \cdot \frac{1}{a}$
- ▶  $\frac{\partial z}{\partial x_1} = \frac{\partial z}{\partial a} \cdot \frac{\partial a}{\partial x_1} = \frac{y_2}{a} \cdot x_2$
- ▶  $\frac{\partial z}{\partial x_2} = \frac{\partial z}{\partial a} \cdot \frac{\partial a}{\partial x_2} + \frac{\partial z}{\partial y_2} \cdot \frac{\partial y_2}{\partial x_2}$   
 $= \frac{y_2}{a} \cdot x_1 + \cos(x_2)$



<sup>9</sup>Source: <https://pytorch.org/blog/overview-of-pytorch-autograd-engine/>

# AutoGrad<sup>10</sup>: Automatic Differentiation in PyTorch

Components of Autograd:

- **tools/autograd**: Contains definitions of derivatives (`derivates.yaml`), and the necessary scripts to compute derivatives. Mostly written in C++.
- **torch/autograd**: Contains Python classes for users to write their own functions (`torch.autograd.Function`) and their gradients (`functional.py`)
- **torch/csrc/autograd**: Contains graph creation files and other execution-related code. Mostly written in C++.

```
>>> x = torch.tensor([0.5, 0.75], requires_grad=True)
```

```
void TensorImpl::set_requires_grad(bool requires_grad) {  
    ...  
    if (!autograd_meta_)  
        autograd_meta_ = impl::GetAutogradMetaFactory()->make();  
    autograd_meta_->set_requires_grad(requires_grad, this);  
}
```

Instantiates this object to hold graph information,  
original class defined in `torch/csrc/autograd/variable.h`

```
struct TORCH_API AutogradMeta : public c10::AutogradMetaInterface {  
    std::string name_;  
  
    Variable grad_;  
    std::shared_ptr<Node> grad_fn_;  
    std::weak_ptr<Node> grad_accumulator_  
    // other fields and methods  
    ...  
};
```

Stores the computed gradient

Pointer used to compute the gradient

Pointer used to add all the gradients wherever this tensor is involved

<sup>10</sup>Source: <https://github.com/pytorch/pytorch>

# Autograd: Computing the Gradient

```
variable_list MulBackward0::apply(variable_list&& grads) {  
    std::lock_guard<std::mutex> lock(mutex_);
```

```
    IndexRangeGenerator gen;  
    auto self_ix = gen.range(1);  
    auto other_ix = gen.range(1);  
    variable_list grad_inputs(gen.size());  
    auto& grad = grads[0];  
    auto self = self_.unpack();  
    auto other = other_.unpack();  
    bool any_grad_defined = any_variable_defined(grads);  
    if (should_compute_output({ other_ix })) {  
        auto grad_result = any_grad_defined ? (mul_tensor_backward(grad, self, other_scalar_type)) : Tensor();  
        copy_range(grad_inputs, other_ix, grad_result);  
    }  
    if (should_compute_output({ self_ix })) {  
        auto grad_result = any_grad_defined ? (mul_tensor_backward(grad, other, self_scalar_type)) : Tensor();  
        copy_range(grad_inputs, self_ix, grad_result);  
    }  
    return grad_inputs;  
}
```

```
struct Edge {  
    ...  
    /// The function this 'Edge' points to.  
    std::shared_ptr<Node> function;  
    /// The identifier of a particular input to the function.  
    uint32_t input_nr;  
};
```

```
>>> x = torch.tensor([0.5, 0.75], requires_grad=True)  
>>> v = x[0] * x[1]  
>>> v  
tensor(0.3750, grad_fn=<MulBackward0>)
```

Generates this function automatically using `tools/autograd` based on the script present in `derivatives.yaml`

Inherited from `TraceableFunction` class, a descendant of `Node` with just a property to enable tracing for debugging/optimize purposes.

```
struct TORCH_API Node : std::enable_shared_from_this<Node> {  
    ...  
    /// Evaluates the function on the given inputs and returns the result of the  
    /// function call.  
    variable_list operator()(variable_list&& inputs) {  
        ...  
    }  
  
protected:  
    /// Performs the 'Node's actual operation.  
    virtual variable_list apply(variable_list&& inputs) = 0;  
    edge_list next_edges_;
```

**But, how is graph constructed, and how are the nodes linked automatically?**

# Autograd: Linking Nodes (Tensor Multiplication)

```
at::Tensor mul_Tensor(c10::DispatchKeySet ks, const at::Tensor & self, const at::Tensor & other) {
    ...
    auto _any_requires_grad = compute_requires_grad( self, other );
    std::shared_ptr<MulBackward0> grad_fn;
    if (!_any_requires_grad) {
        // Creates the link to the actual grad_fn and links the graph for backward traversal
        grad_fn = std::shared_ptr<MulBackward0>(new MulBackward0(), deleteNode);
        grad_fn->set_next_edges(collect_next_edges( self, other ));
    }
    ...
    // Does the actual function call to ATen
    auto tmp = (c10::) {
        at::AutoDispatchBelowADInplaceOrView guard;
        return at::redispatch::mul(ks & c10::after_autograd_keyset, self_, other_);
    }();

    auto result = std::move(tmp);
    if (grad_fn) {
        // Connects the result to the graph
        set_history(flatten_tensor_args( result ), grad_fn);
    }
    ...
    return result;
}
```

Creates the **grad\_fn** object and the edges needed to link the nodes

Connects all output tensors to **grad\_fn** node after forward pass.

```
void set_next_edges(edge_list&& next_edges) {
    next_edges_ = std::move(next_edges);
    for(const auto& next_edge : next_edges_) {
        update_topological_nr(next_edge);
    }
}
```

```
struct MakeNextFunctionList : IterArgs<MakeNextFunctionList> {
    edge_list next_edges;
    using IterArgs<MakeNextFunctionList>::operator();
    void operator()(const Variable& variable) {
        if (variable.defined()) {
            next_edges.push_back(input_gradient_edge(variable));
        } else {
            next_edges.emplace_back();
        }
    }
    void operator()(const c10::optional<Variable&& variable) {
        if (variable.has_value() && variable.defined()) {
            next_edges.push_back(input_gradient_edge(variable));
        } else {
            next_edges.emplace_back();
        }
    }
};

template <typename... Variables>
edge_list collect_next_edges(Variables&&... variables) {
    detail::MakeNextFunctionList make;
    make.apply(std::forward<Variables>(variables)...);
    return std::move(make.next_edges);
}
```

Creating an **Edge** object per input variable

```
void set_gradient_edge(const Variable& self, Edge edge) {
    auto meta = materialize_autograd_meta(self);
    meta->grad_fn_ = std::move(edge.function);
    meta->output_nr_ = edge.input_nr;
    // For views, make sure this new grad_fn_ is not overwritten unless it is necessary
    // in the VariableHooks::grad_fn below.
    // This logic is only relevant for custom autograd Functions for which multiple
    // operations can happen on a given Tensor before its gradient edge is set when
    // exiting the custom Function.
    auto diff_view_meta = get_view_autograd_meta(self);
    if (diff_view_meta && diff_view_meta->has_base_view()) {
        diff_view_meta->set_attr_version(self._version());
    }
}
```

Copies everything to **Autogradmeta** object

```
inline void set_history(
    at::Tensor& variable,
    const std::shared_ptr<Node0&& grad_fn) {
    AT_ASSERT(grad_fn);
    if (variable.defined()) {
        // If the codegen triggers this, you most likely want to add your newly added function
        // to the DONT_REQUIRE_DERIVATIVE list in tools/autograd/gen_variable_type.py
        TORCH_INTERNAL_ASSERT(isDifferentiableType(variable.scalar_type()));
        auto output_nr =
            grad_fn->add_input_metadata(variable);
        input_set_gradient_edge(variable, {grad_fn, output_nr});
    } else {
        grad_fn->add_input_metadata(Node::undefined_input());
    }
}
```

```
Edge gradient_edge(const Variable& self) {
    // If grad_fn is null (as is the case for a leaf node), we instead
    // interpret the gradient function to be a gradient accumulator, which will
    // accumulate its inputs into the grad property of the variable. These
    // nodes get suppressed in some situations, see "suppress gradient
    // accumulation" below. Note that only variables which have "requires_grad =
    // True" can have gradient accumulators.
    if (const auto& gradient = self.grad_fn()) {
        return Edge(gradient, self.output_nr());
    } else {
        return Edge(grad_accumulator(self), 0);
    }
}
```