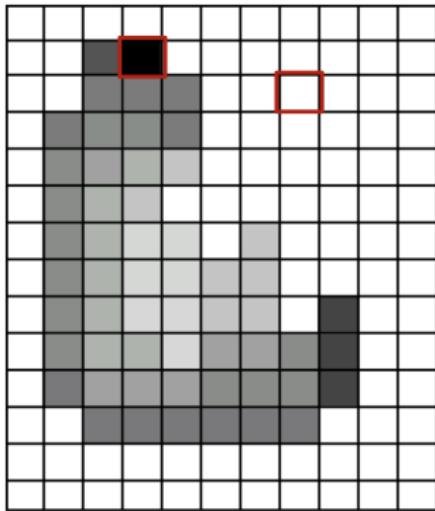


Topic 3: Deep Learning in Computer Vision

Image Representation: Matrices

- **Matrix:** Grayscale images



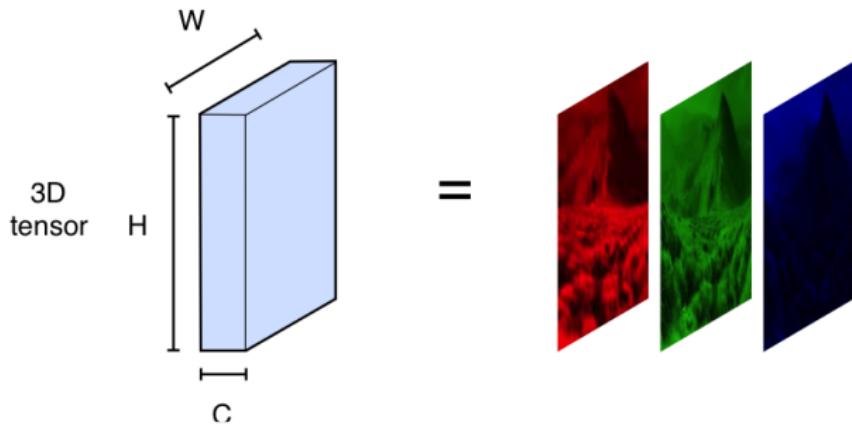
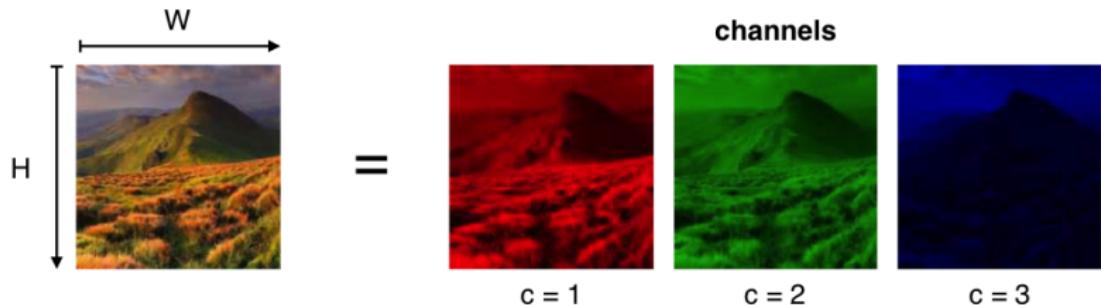
=

255	255	255	255	255	255	255	255	255	255	255	255
255	255	20	0	255	255	255	255	255	255	255	255
255	255	75	75	255	255	255	255	255	255	255	255
255	75	95	95	75	255	255	255	255	255	255	255
255	96	127	145	175	255	255	255	255	255	255	255
255	127	145	175	175	175	255	255	255	255	255	255
255	127	145	200	200	175	175	95	255	255	255	255
255	127	145	200	200	175	175	95	47	255	255	255
255	127	145	145	175	127	127	95	47	255	255	255
255	74	127	127	127	95	95	95	47	255	255	255
255	255	74	74	74	74	74	74	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255

0 = black; 255 = white

Image Representation: Tensors

- **Tensor:** Most common form of representation

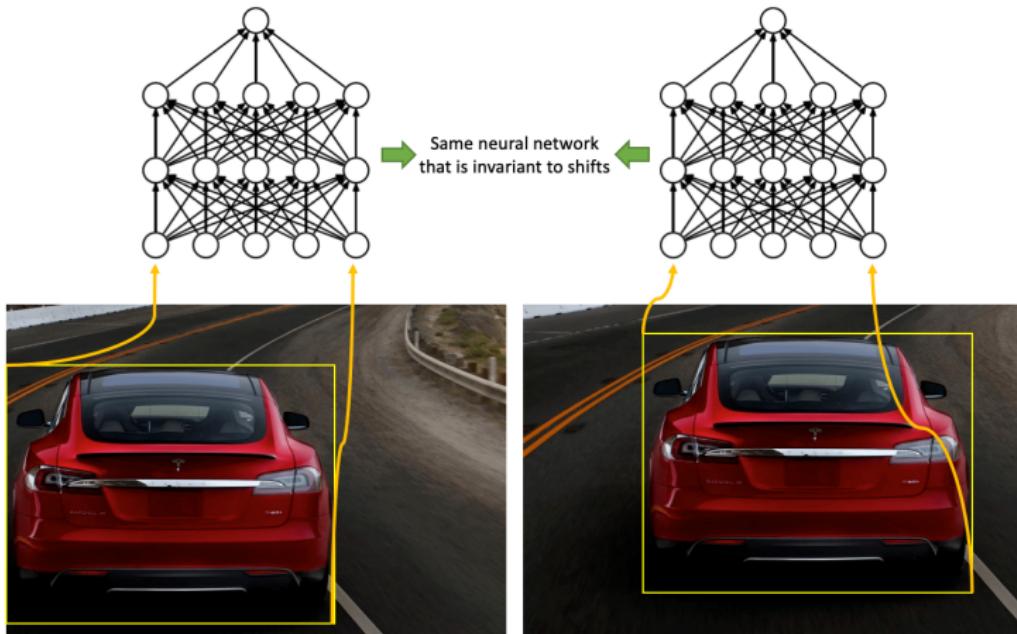


Computer Vision: An Collection of Exciting/Challenging Problems

- ▶ Object Detection
- ▶ Object Localization
- ▶ Object Tracking
- ▶ Segmentation
- ▶ Alignment

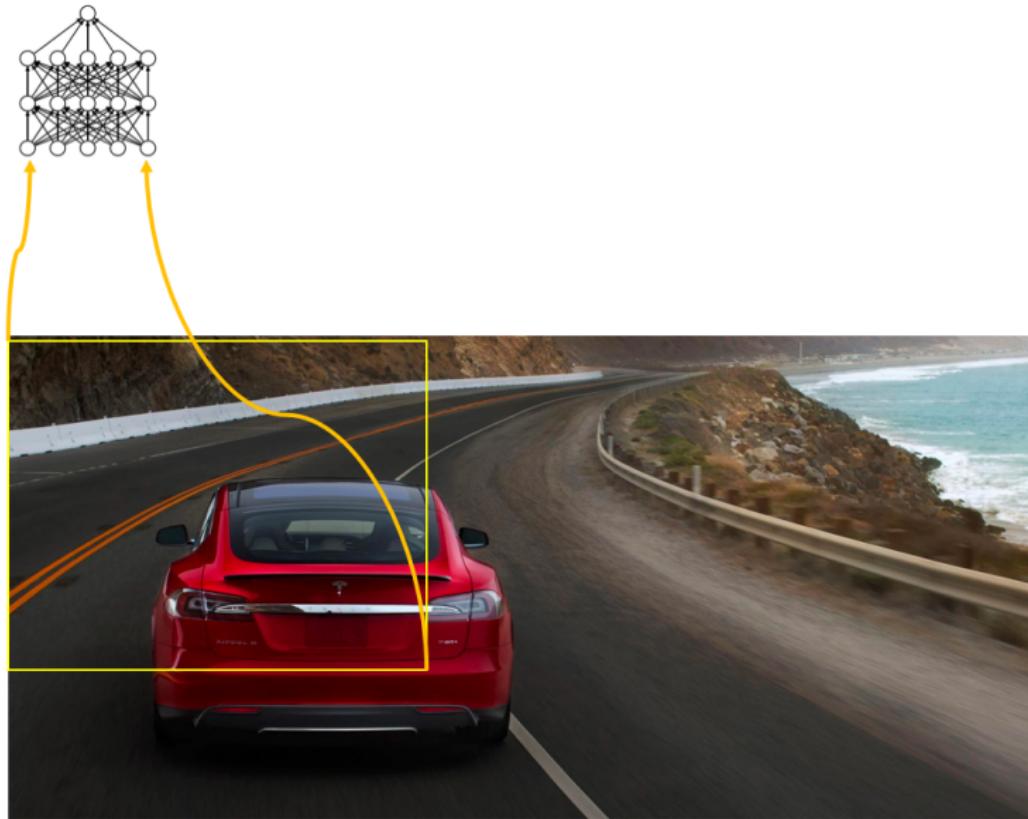
Object Detection with NNs: Need for Shift Invariance

- ▶ Will a neural network trained to detect a car in the left corner of the image, detect a car in a different location?
 - ▶ Conventional neural networks are sensitive to *shifts*.
 - ▶ Need one that detects the presence of car regardless of its location.



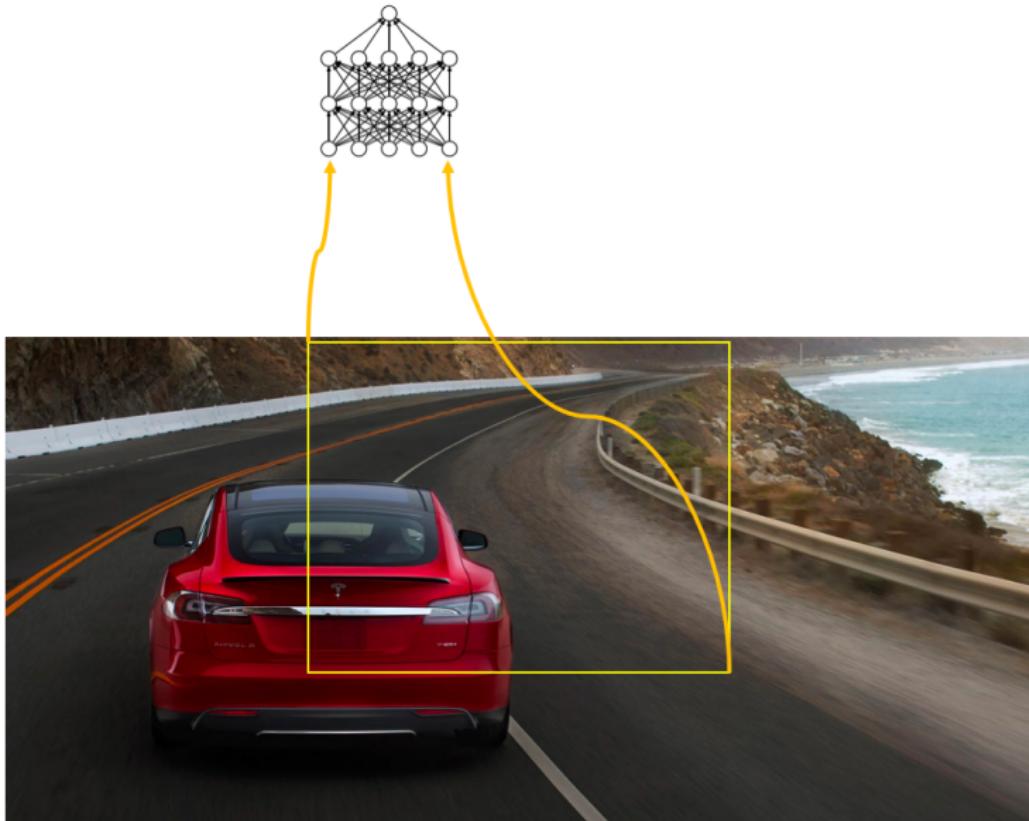
Solution: Scan the Image

- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.



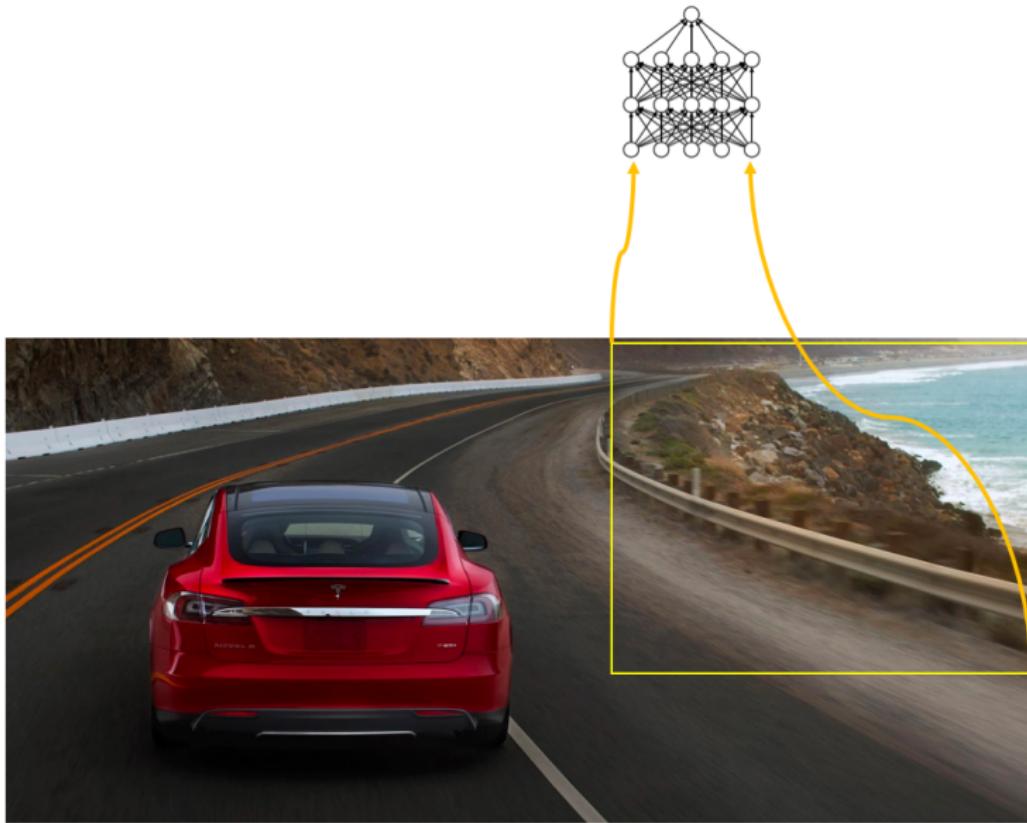
Solution: Scan the Image

- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.



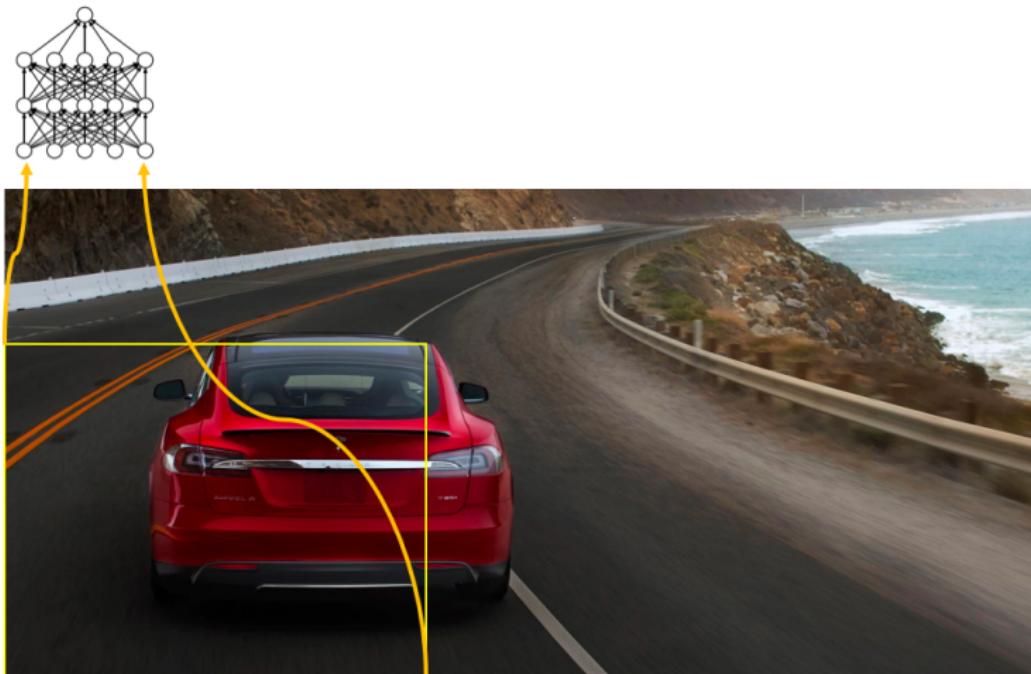
Solution: Scan the Image

- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.



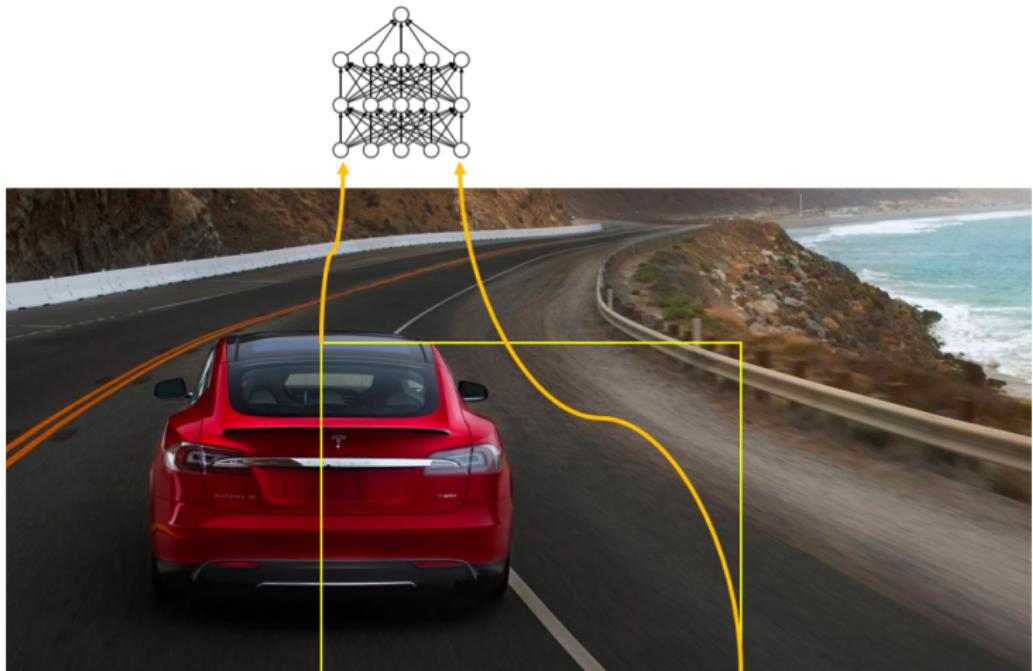
Solution: Scan the Image

- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.



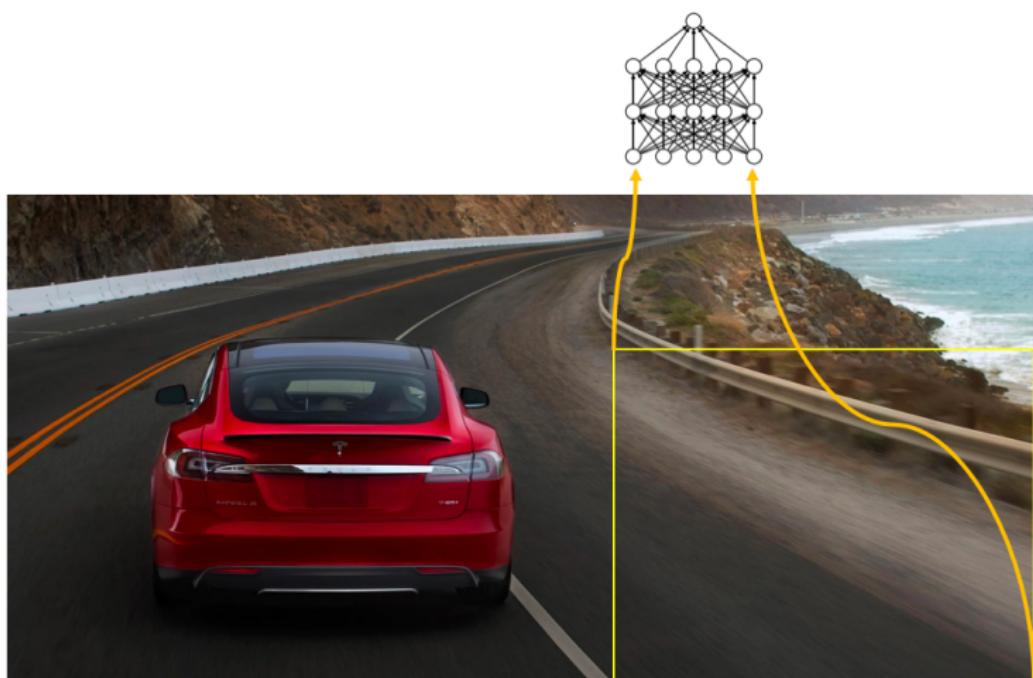
Solution: Scan the Image

- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.

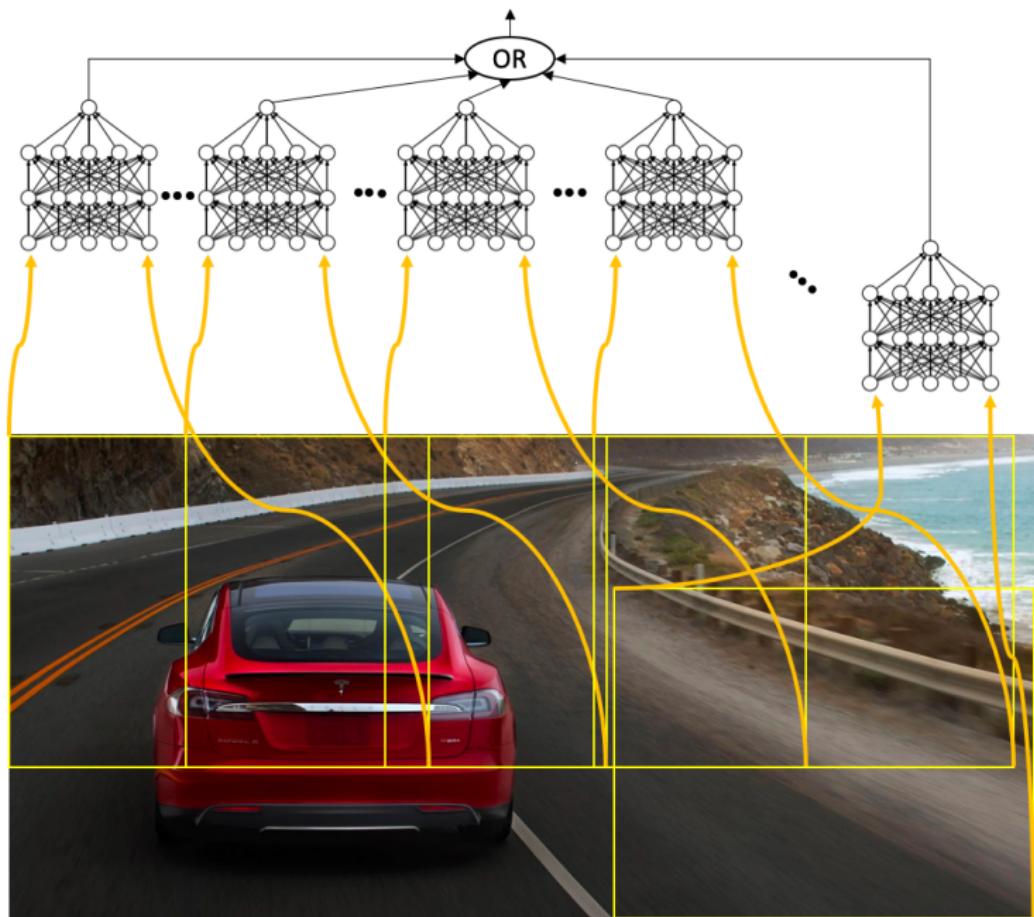


Solution: Scan the Image

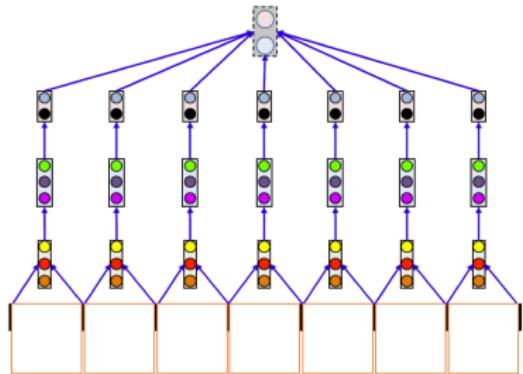
- ▶ Scan for the target object at each location using a patch (kernel) of size $K \times K$.



Combining all the scans, we get a Giant Neural Network!



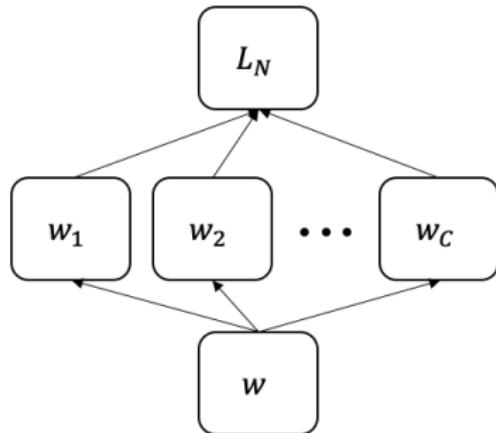
Note: This is a Shared Parameter Network!



```
def giant_nn(Img,K):
    W = width(Img)
    H = height(Img)
    C = num_channels(Img)
    for i = 1:W-K+1
        for j = 1:H-K+1
            ImgSegment = Img(1:C,i:i+K-1, j:j+K-1)
            y(i,j) = NN(ImgSegment)
    Y = OR(y(1,1), ..., y(W-K+1,H-K+1))
```

- ▶ Note that this is not the same as regular NN models.
- ▶ **Shared Parameter Networks:**
 - ▶ `NN()` uses the same weights for different locations
 - ▶ Block-structured weight matrix, with identical blocks for each kernel scan!
 - ▶ In other words, any parameter update of one scan must update equally for all scans.

Training Shared Parameter Networks

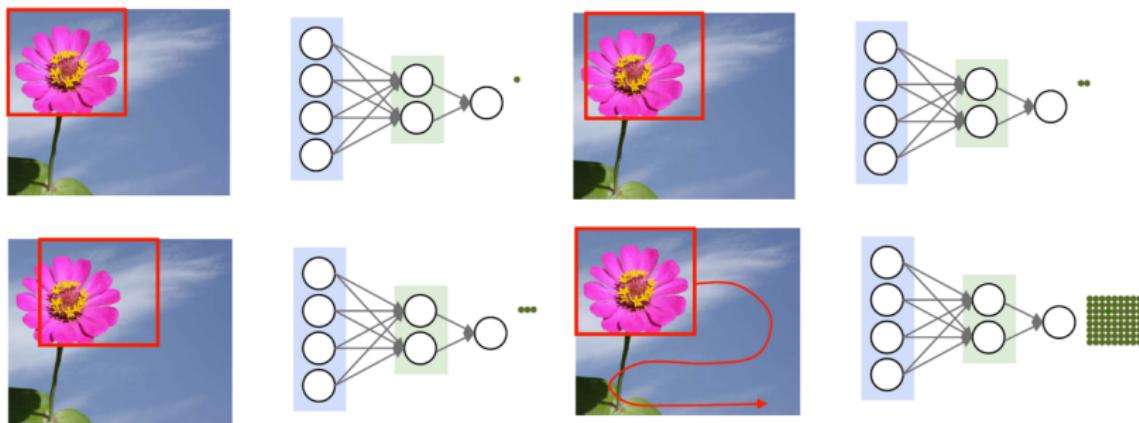


$$\begin{aligned}\frac{\partial L_N}{\partial w} &= \sum_{c=1}^C \frac{\partial L_N}{\partial w_c} \cdot \frac{\partial w_c}{\partial w} \\ &= \sum_{c=1}^C \frac{\partial L_N}{\partial w_c}.\end{aligned}$$

- ▶ Say, we have L layers in $\text{NN}(\cdot)$ \Rightarrow Let the weights be W_1, \dots, W_L
- ▶ Let \mathcal{S} denote the set of all edges with common value
- ▶ $\nabla_{w_{\mathcal{S}}} L_N = \frac{\partial L_N}{\partial w_{\mathcal{S}}} = \sum_{c \in \mathcal{S}} \frac{\partial L_N}{\partial w_c}.$
- ▶ Update step: $w_{\mathcal{S}} \leftarrow w_{\mathcal{S}} - \eta \nabla_{w_{\mathcal{S}}} L_N,$
- ▶ For every $(\ell, i, j) \in \mathcal{S}$, define $w_{i,j}^{\ell} = w_{\mathcal{S}}.$
- ▶ Continue until the stopping criterion is satisfied.

Take a Closer Look at Scanning...

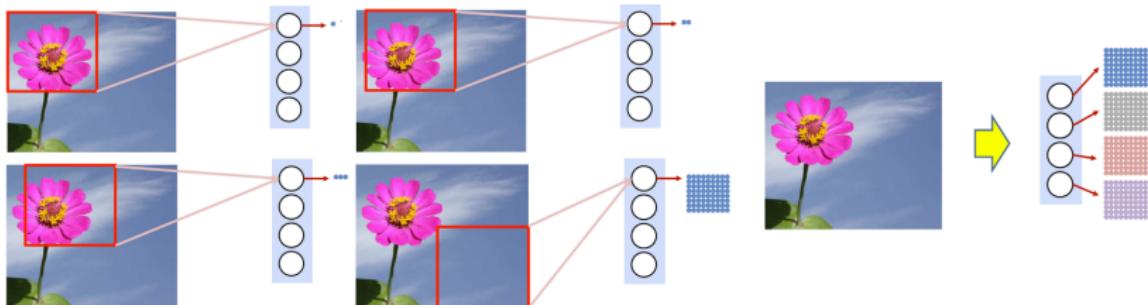
- ▶ Each scan passes the image segment through the entire network and produces an outcome
- ▶ Combining all the scans, we get a collection of outcomes, which are passed through OR (or softmax, or any other flat-NN) gate.



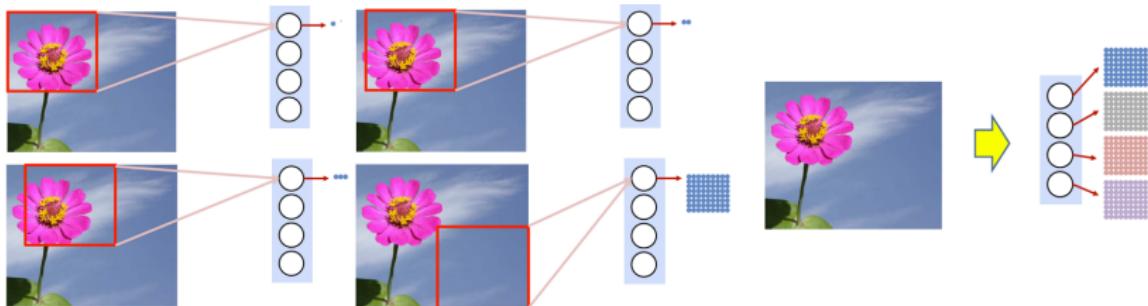
What if, we compute the outcome of one neuron for all scans?

A New Scanning Approach...

- ▶ Scan the entire image with one neuron
- ▶ Arrange the neuron's outputs from the scanned positions according to their positions in the original image

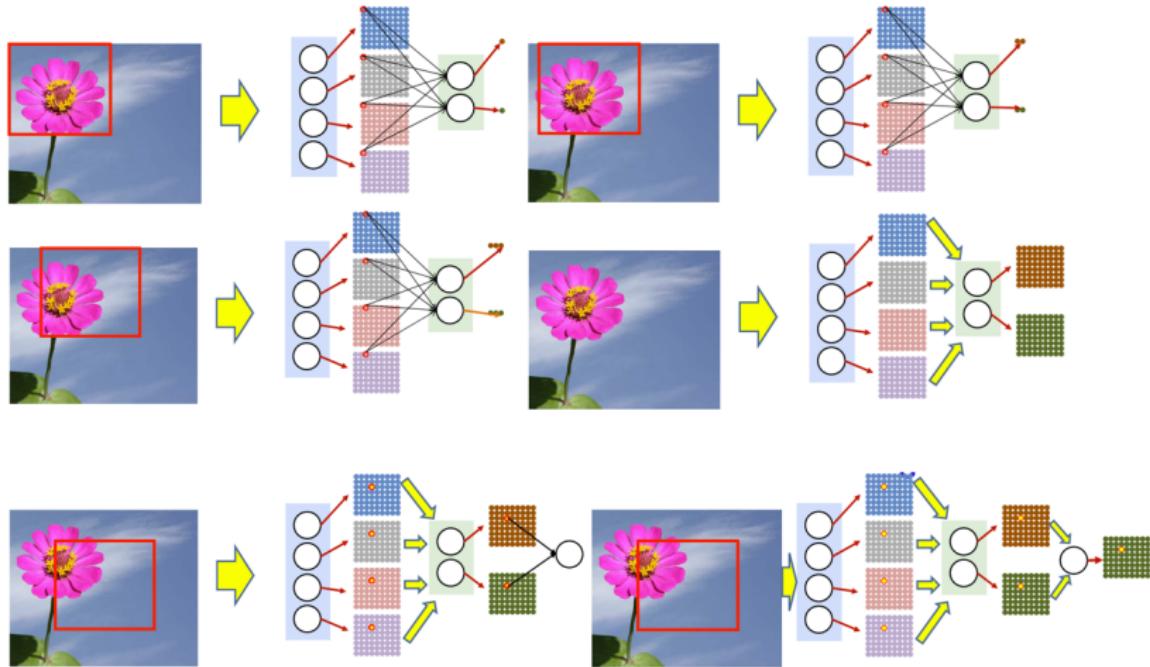


- ▶ Recurse the same logic across layers...



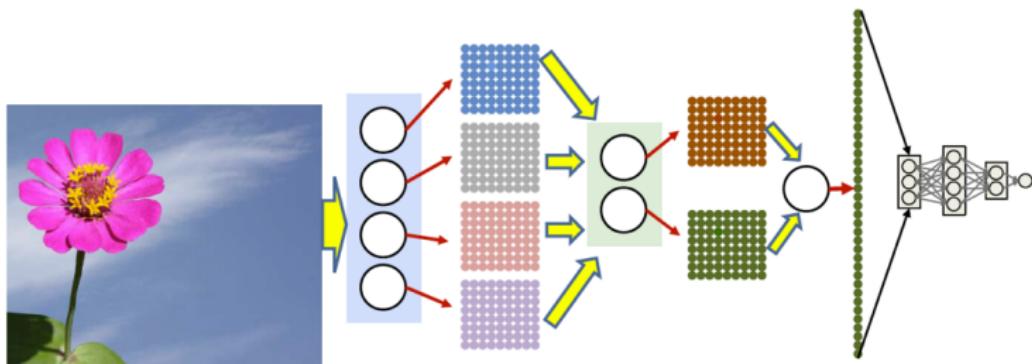
A New Scanning Approach (cont...)

- ▶ Subsequent layers jointly scan multiple maps of the previous layer.
- ▶ Final layer ⇒ A map that detects flower at each location in the image.

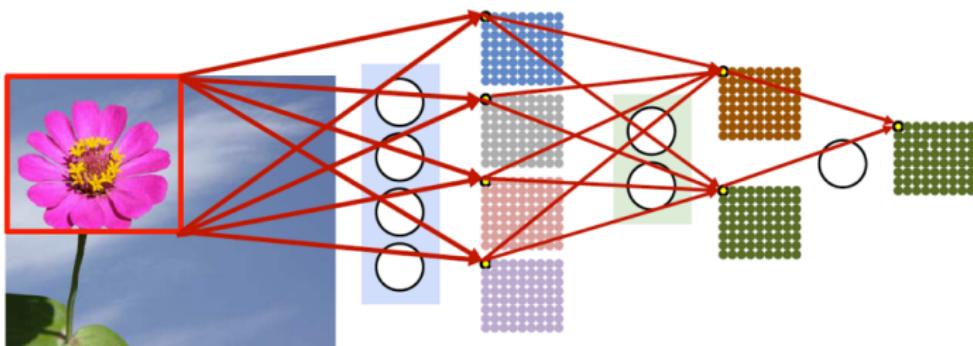


A New Scanning Approach (cont...)

- Goal: Detect the flower, not locate it!
- Flatten the final map and pass it into a softmax layer, or another NN.

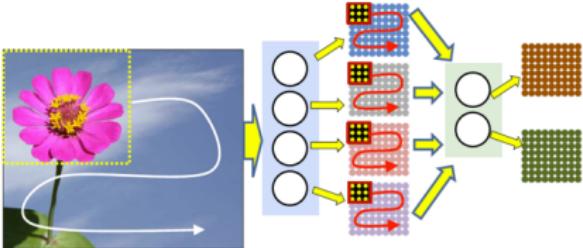
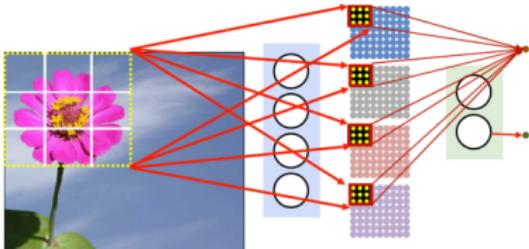
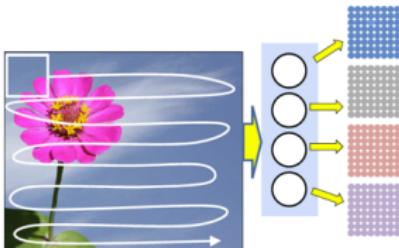


Is this scan sufficient?



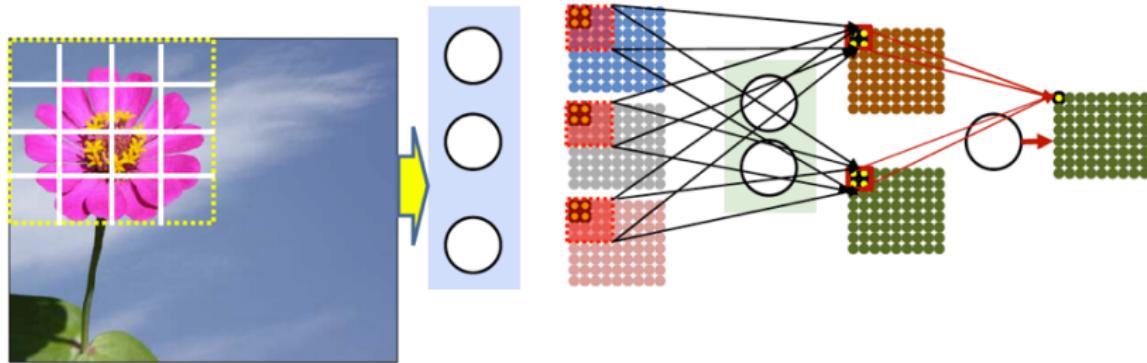
A New Scanning Approach (cont...)

- ▶ Distribute the scan \Rightarrow Subsequent layers evaluate blocks of outputs from previous layers
- ▶ Now, we can evaluate larger windows
 - ▶ Windows distributed over layers \Rightarrow Useful to learn larger patterns
- ▶ Jointly scan all first-layer maps \Rightarrow Output of second layer represents a scan of full-sized input window



Convolutional Neural Network (CNNs)

Use the same logic recursively across multiple layers,
we obtain a convolutional neural network!



```
def cnn(Img, L, K):
    W = width(Img)
    H = height(Img)
    C = num_channels(Img)
    for l = 1:L
        Compute  $W_{l-1}, H_{l-1}, K_l, D_l, D_{l-1}$ 
        for x = 1: $W_{l-1} - K_l + 1$ 
            for y = 1: $H_{l-1} - K_l + 1$ 
                Segment = Y(1-1, 1: $D_{l-1}$ , x:x+ $K_l-1$ , y:y+ $K_l-1$ )
                for j = 1: $D_l$ 
                    Compute z(l, j, x, y) from Segment
                    Y(l, j, x, y) = activation(z(l, j, x, y))
    Y = softmax(Y(L, :, 1, 1), ..., Y(L, :, W+K-1, H-K+1))
```